

## 2교시.

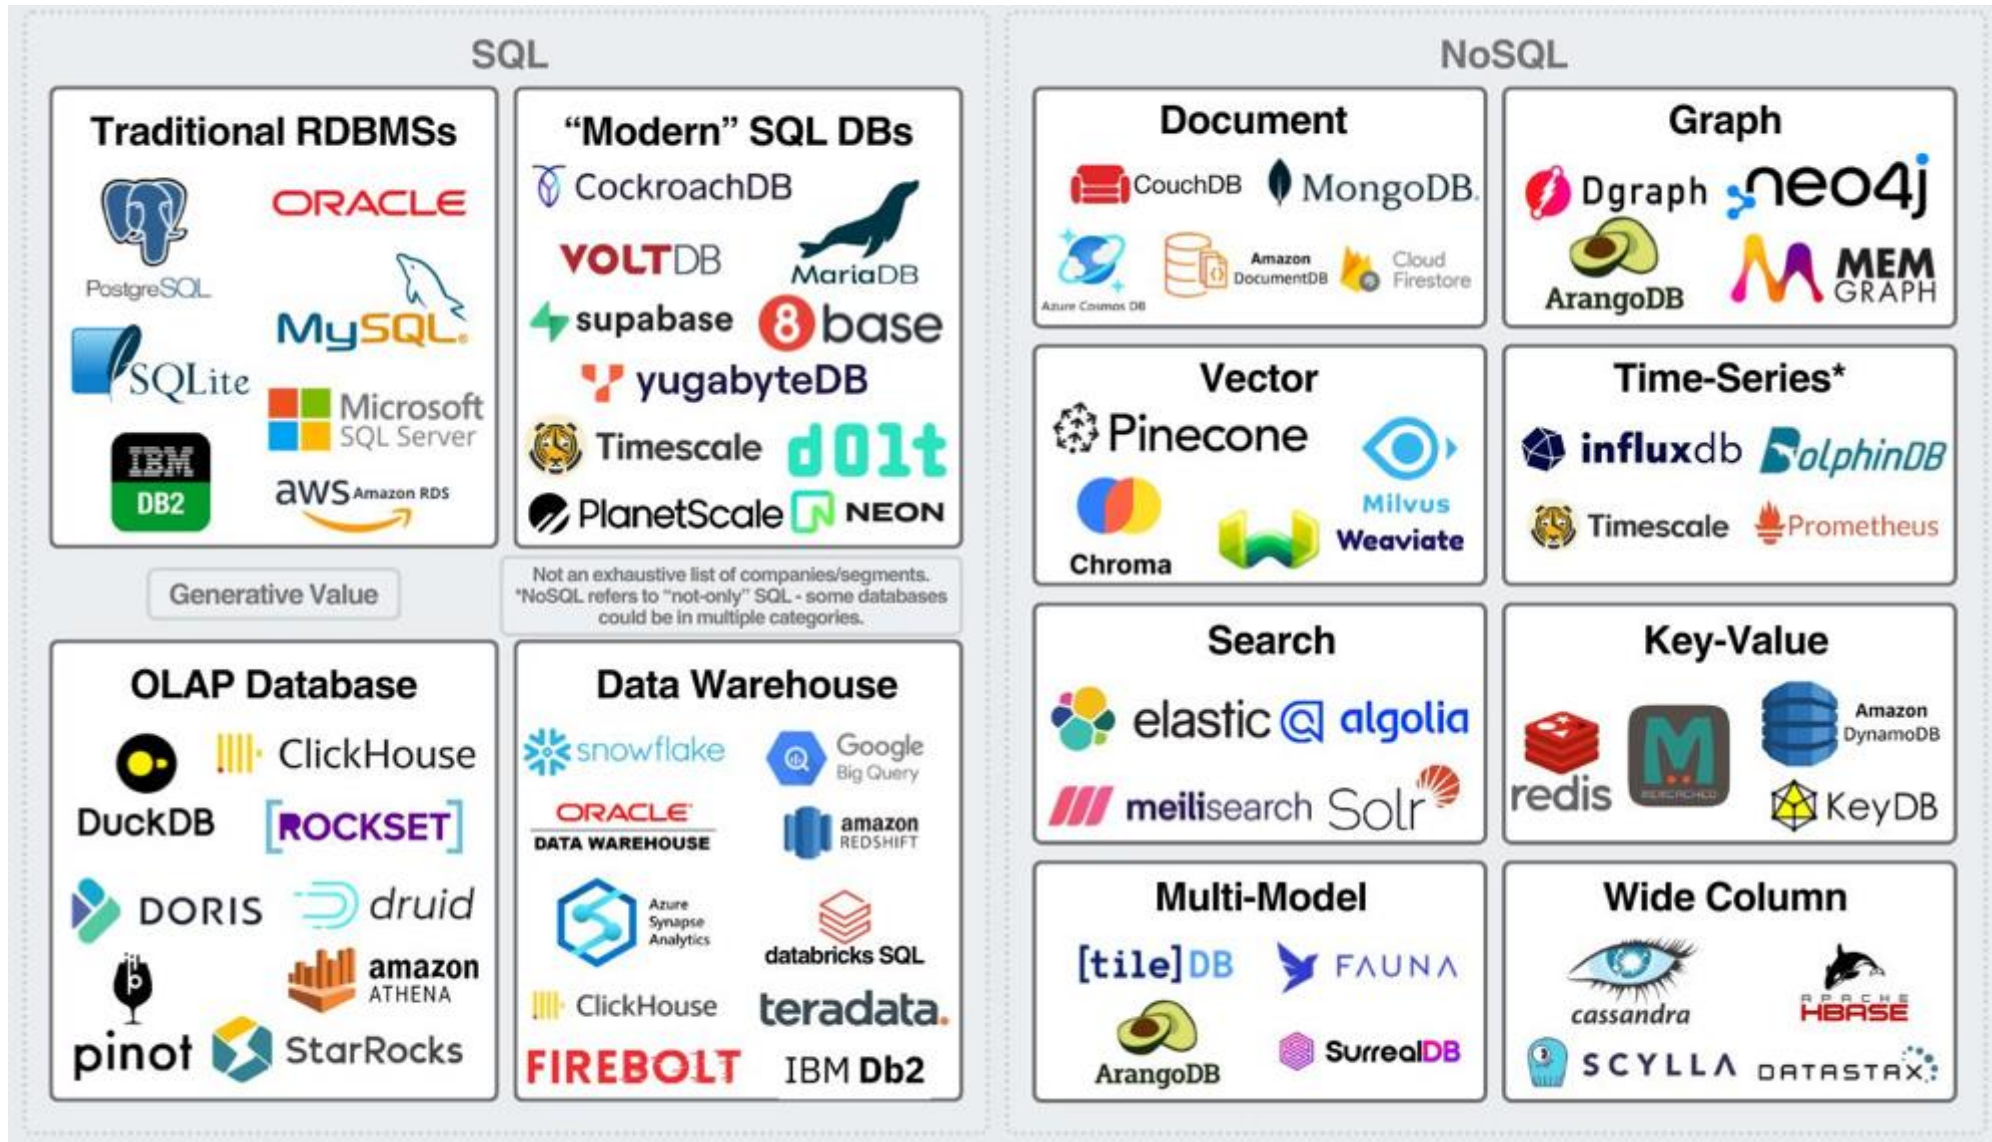
# NoSQL 모델링, SDK 개요, 실습

- 1 SQL, NoSQL 모델링
- 2 SDK 개요
- 3 SDK 실습 : Python, Java
- A SQL Samples

## 2-1. SQL, NoSQL 모델링



# 데이터 모델 : SQL(Table)과 NoSQL(Real World)



# 데이터 모델 : JSON 도큐먼트

JSON은 텍스트로 이루어져 있으므로, 사람과 기계 모두 읽고 쓰기 쉽다. 프로그래밍 언어와 플랫폼에 독립적이므로, 서로 다른 시스템 간에 객체를 교환하기에 좋다.

- JSON 도큐먼트의 장점
  - 단일 도큐먼트 내에 다양한 정보를 계층 구조를 활용하여 저장
  - 정보 추가/삭제가 유연한 구조 제공
  - 데이터 전달을 위한 표준 인터페이스 역할
- RDB와 차별점
  - 여러 테이블로 분리, 저장되는 데이터를 단일 도큐먼트에 저장
  - 테이블 간 조인을 최소화하여 데이터 처리 속도 향상

SQL, Table

이 력 서

1. 기초자료

성명	홍길동	성별	남	생년월일	1999-01-01
주민등록번호	000000-000000000				
E-mail	bizforms@bizforms.co.kr				
전화번호	02-000-0000	휴대폰	000-123-4567		
우편번호	00000	주소	00번지 00번지		
주거형태	00번지	관계	장남		
호적관계	호주				

기초정보

2. 학력사항

년/월/일	학	교	명	비	고
2000-02	00정보	공업고등학교	졸업		

학력사항

경력사항

개인능력

000.00	담당
--------	----

관련내용

리눅스(레드햇기반)운영	비고
samba서버 운영(nfs기반)	
apache웹서버 운영	
sendmail서버 운영	
shell 스크립트(bash) 작성 가능	
perl 및 C도 기본적으로 이해 가능	
DNS서버 운영	
WINDOWS2000 서버 운영 및 BS, AD(분산파일시스템), DNS	
NAT, DHCP, SMTP, POP3 등의 service 운영가능	
전반적인 네트워크구성 이해(TCP/IP기반)	

하드웨어에 맞춘 데이터 모델

NoSQL, JSON 도큐먼트

KEY : 1001

```
{
  "성명": "홍길동",
  "주소": "서울시 00구 00동 000-000",
  "E-mail": "HongKildong@couchbase.com",
  ...
  "학력사항": [
    {
      "졸업년도": "2019년",
      "학교명": "00정보 공업고등학교"
    }
  ],
  "경력사항": [
    {
      "기간": "2019 ~ 현재",
      "관련내용": "XX글로벌 IT팀 Unix 서버 담당"
    }
  ],
  "개인능력": [
    {
      "관련내용": "리눅스 운영"
    },
    {
      "비고": NULL
    }
  ],
  ...
}
```

실제 세계에 맞춘 데이터 모델

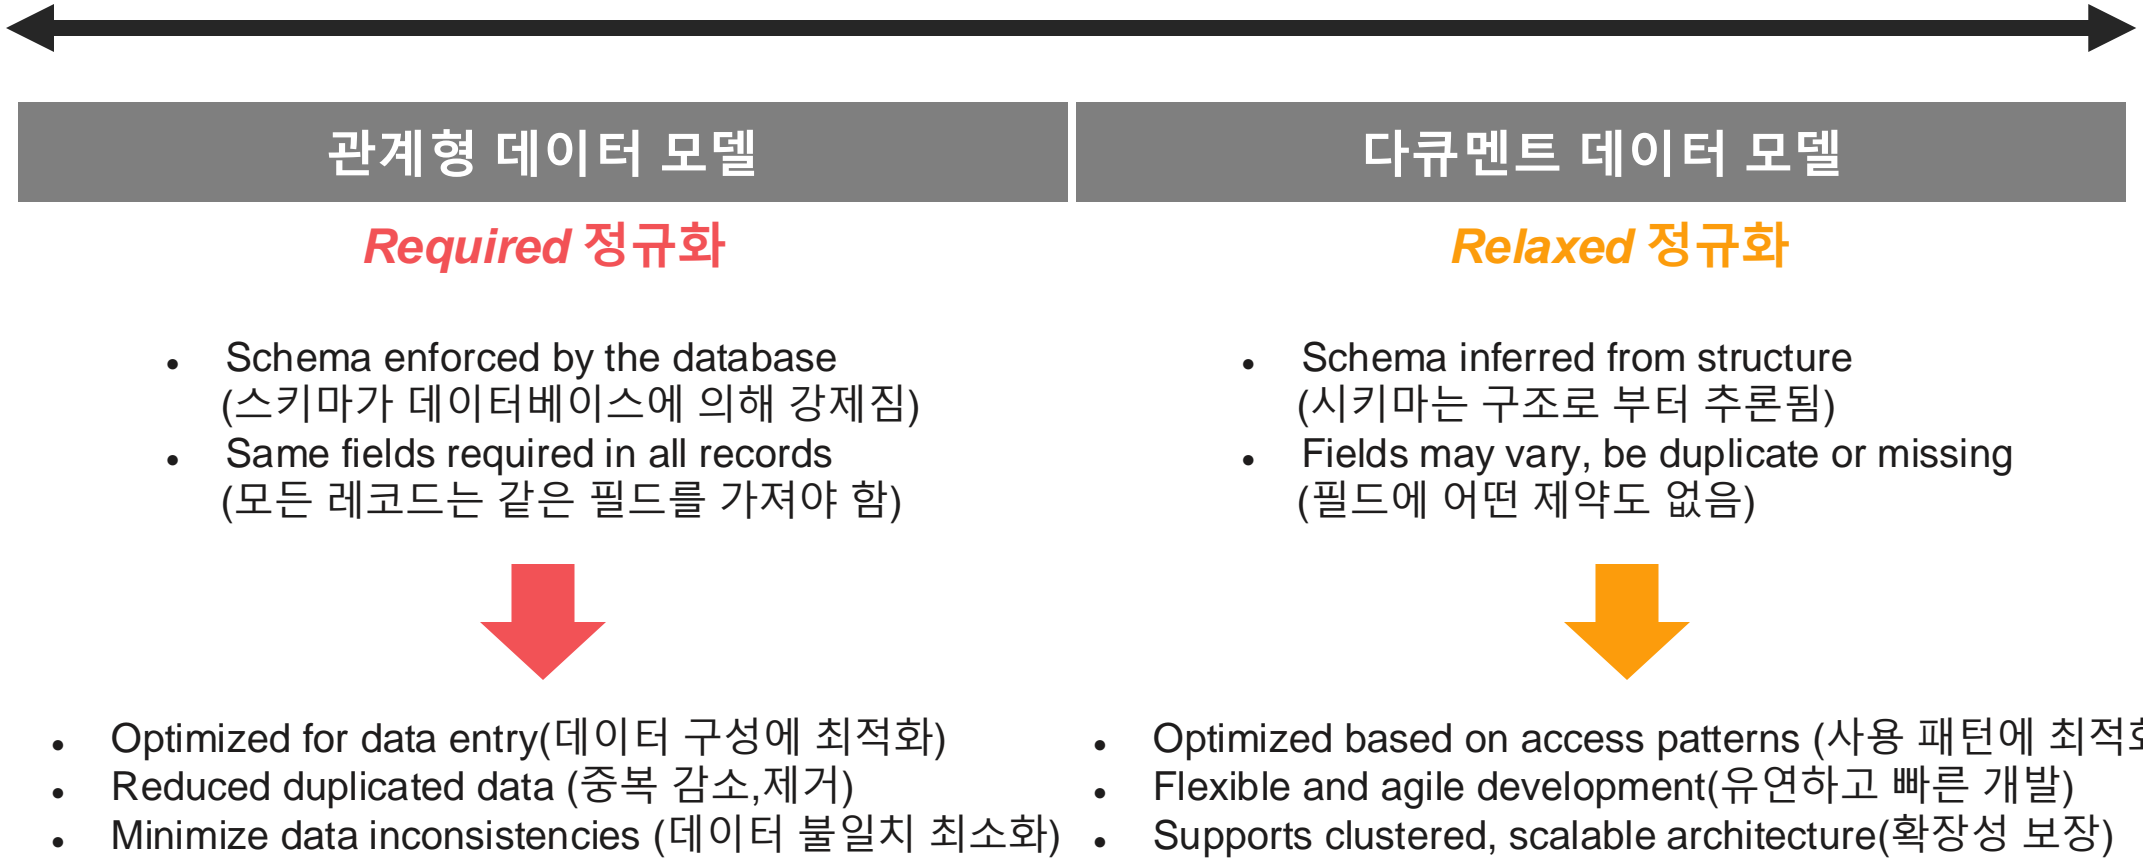
# 논리 / 물리 모델

- RDBMS와 유사한 구조의 논리 계층 구조로 구성하여 편리한 데이터 관리
- Data 서비스를 완전 메모리DB로 사용도 가능하며 용도에 따라 물리 저장 방식을 선택할 수 있음

RDBMS	Couchbase
Server	Cluster
Database	Bucket
Schema	Scope
Table	Collection
Row	Document (JSON)
Value	Sub-Document, Array

Feature	Ephemeral Bucket	Couchbase Bucket	Magma Bucket
Bucket memory quota (per node)	Min 256MB	Min 256MB	Min 1024MB
Max Object Size	20MB	20MB	20MB
Persistence	no	yes	yes
Replication and XDCR	yes	yes	yes
Encrypted data access	yes	yes	yes
Rebalance	yes	yes	yes
N1QL, Seach, Analytics, Eventing	yes	yes	yes
Indexing	yes	yes	yes
Backup	yes	yes	yes

# 관계형 vs. 다큐멘터 데이터 모델



# Real World Data 를 JSON으로 표현

Customer		
CustomerID	Name	DoB
CBL2017	Jane Smith	1990-01-30

Customer Document Key: CBL2017

```
{  
  "Name" : "Jane Smith",  
  "DOB"  : "1990-01-30"  
}
```

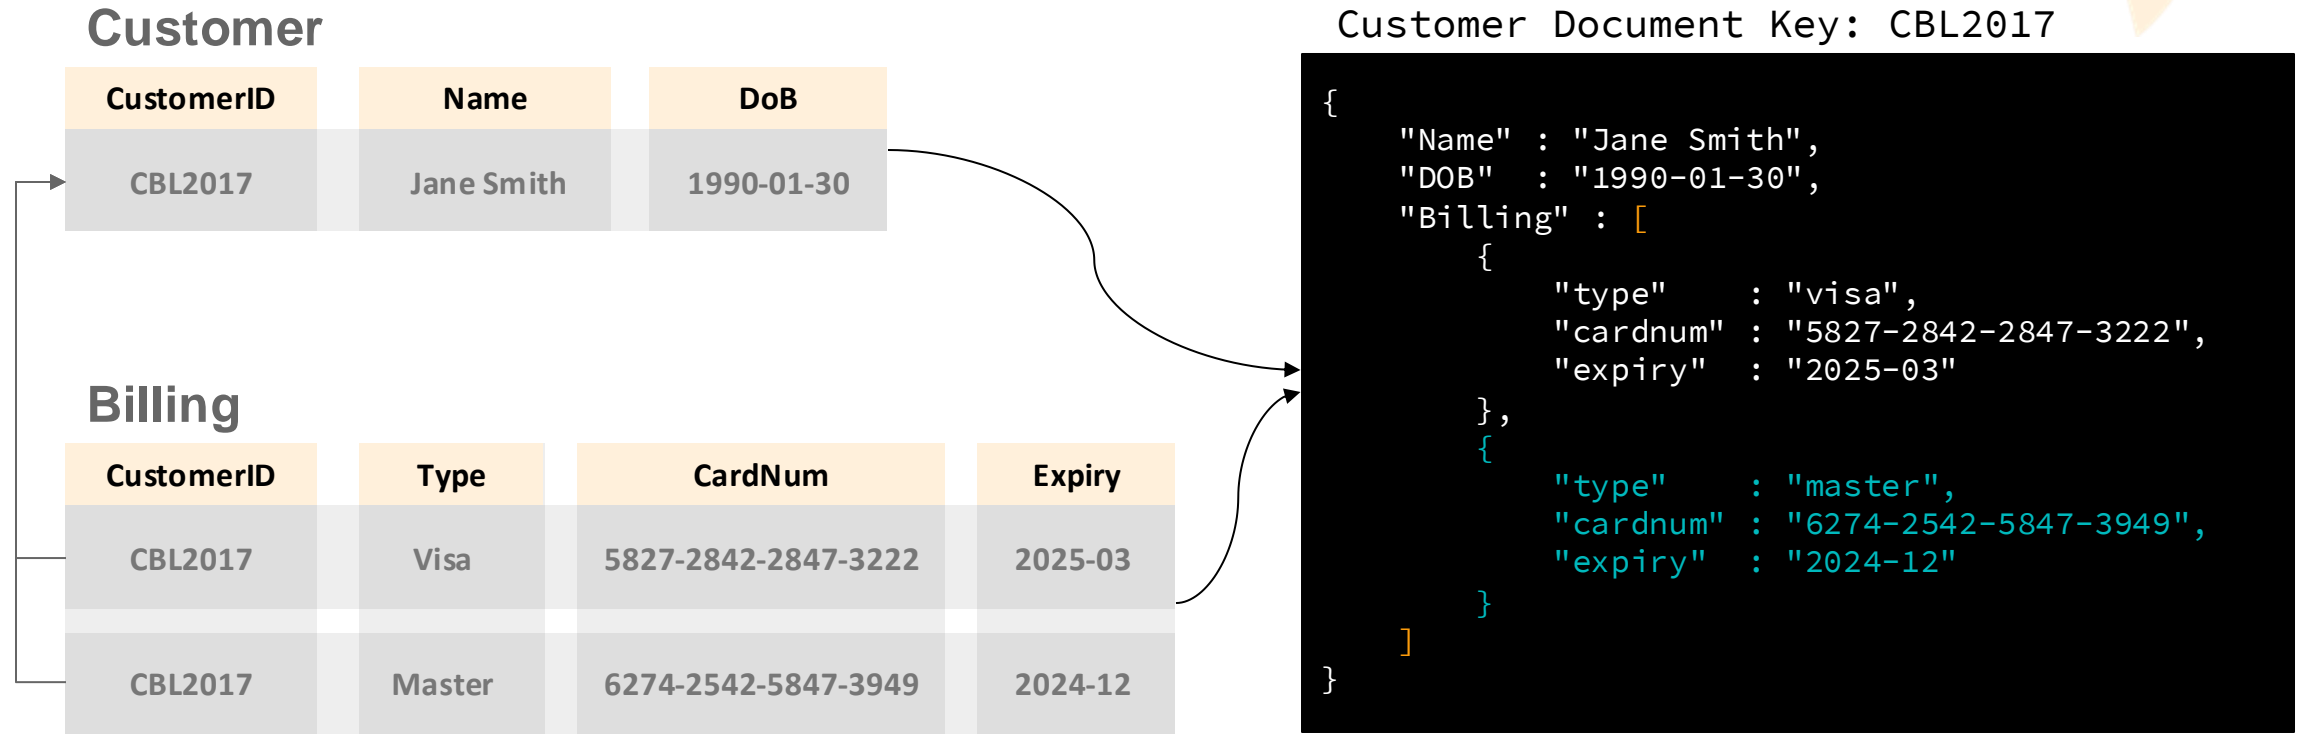
OR

Customer Document Key: CBL2017

```
{  
  "Name" : {  
    "fname": "Jane",  
    "lname": "Smith"  
  },  
  "DOB"  : "1990-01-30"  
}
```

- The Primary Key becomes the Document Key
- Column name-Column value become KEY-VALUE pair

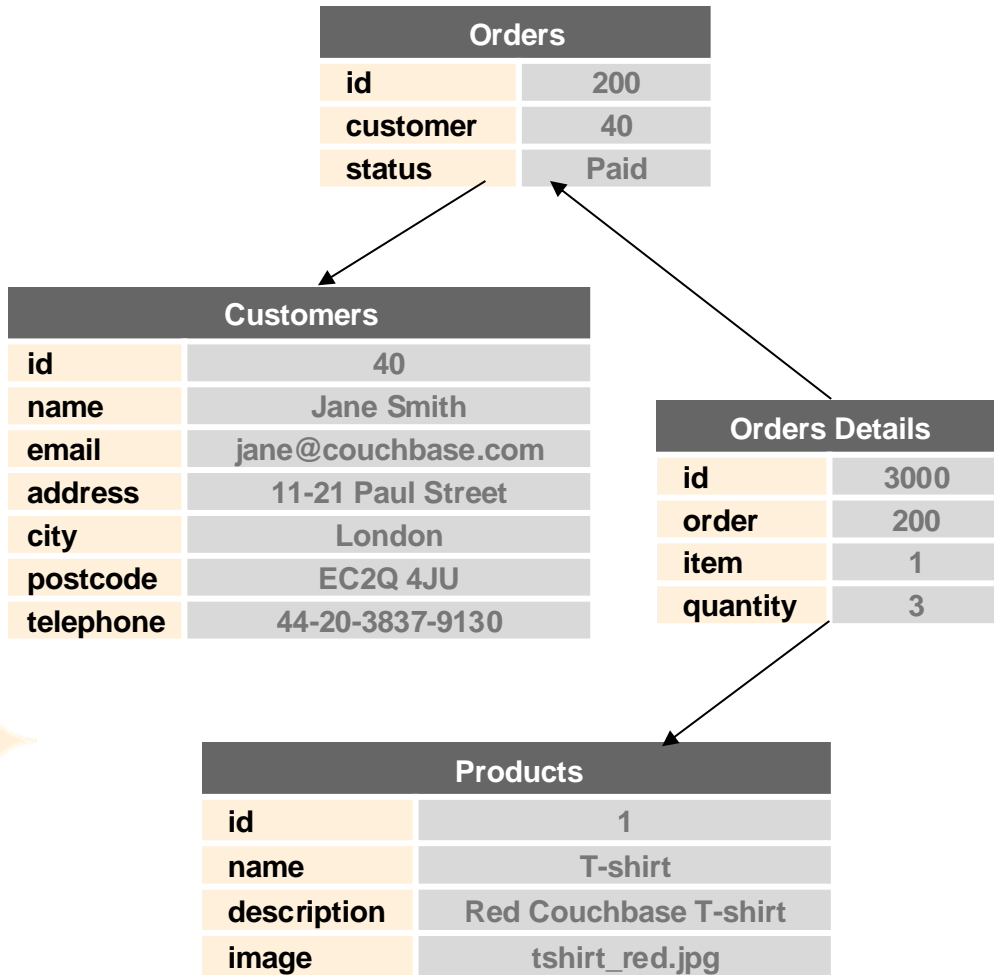
# 고객의 카드 정보를 JSON으로 표현



- Denormalization simplifies data access and offers the best performance  
비정규화는 데이터 사용을 단순화 해주고, 최고의 성능을 제공함
- Value evolution: Simply add additional array element or update a value



# 테이블(RDBMS)을 컬렉션(NoSQL)에 매핑 | eCommerce 예제



Bucket: ecommerce | Scope: \_default

Collection:  
'Products'

```
{
  "name": "T-shirt",
  "description": "Red Couchbase T-shirt",
  "image": "tshirt_red.jpg"
}
```

Doc ID: 1

Collection:  
'Customers'

```
{
  "name": "Jane Smith",
  "email": "jane@couchbase.com",
  "address": "11-21 Paul Street",
  "city": "London",
  "postCode": "EC2A 4JU",
  "telephone": "44-20-3837-9130"
}
```


Doc ID: 40

Collection:  
'Orders'

```
{
  "customer": {
    "id": 40,
    "name": "Jane smith",
    "email": "jane@couchbase.com"
  },
  "status": "Paid",
  "orderDetails": [
    { "productId": 1, "name": "T-shirt", "quantity": 3 },
  ]
}
```


Doc ID: 200

# JSON 데이터 모델시 고려사항



## EMBED WHEN

- There is an Ownership Relationship  
오너쉽 관계가 명확할 때
- Both docs are frequently accessed together  
두 문서가 동시에 사용되는 경우가 많을 때
- Reads greatly outnumber writes  
읽기가 쓰기에 비해 많을 때
- Data is small  
데이터 사이즈가 작을 때



## REFER WHEN

- There is not an Ownership Relationship  
오너쉽 관계가 명확하지 않을 때
- Both docs are not frequently accessed together  
두 문서가 동시에 사용되는 경우가 작을 때
- Document is updated frequently  
문서가 자주 변경될 때
- Need to reduce the document size  
문서의 크기를 줄일 필요가 있을 때

Try to embed first, refer when it makes sense  
먼저 Embed를 시도해 보고, 그 다음 Refer 고려

# Relationships에 따른 Embed 와 Refer

- 1-1

**Embed** Example: Satellite and Manufacturer


- 1-Many

**Embed** Example: Satellite and Missions


**Reference** Example: Measurements, Satellites and Instruments

- Many-Many

**Reference** Example: Satellite and Instruments

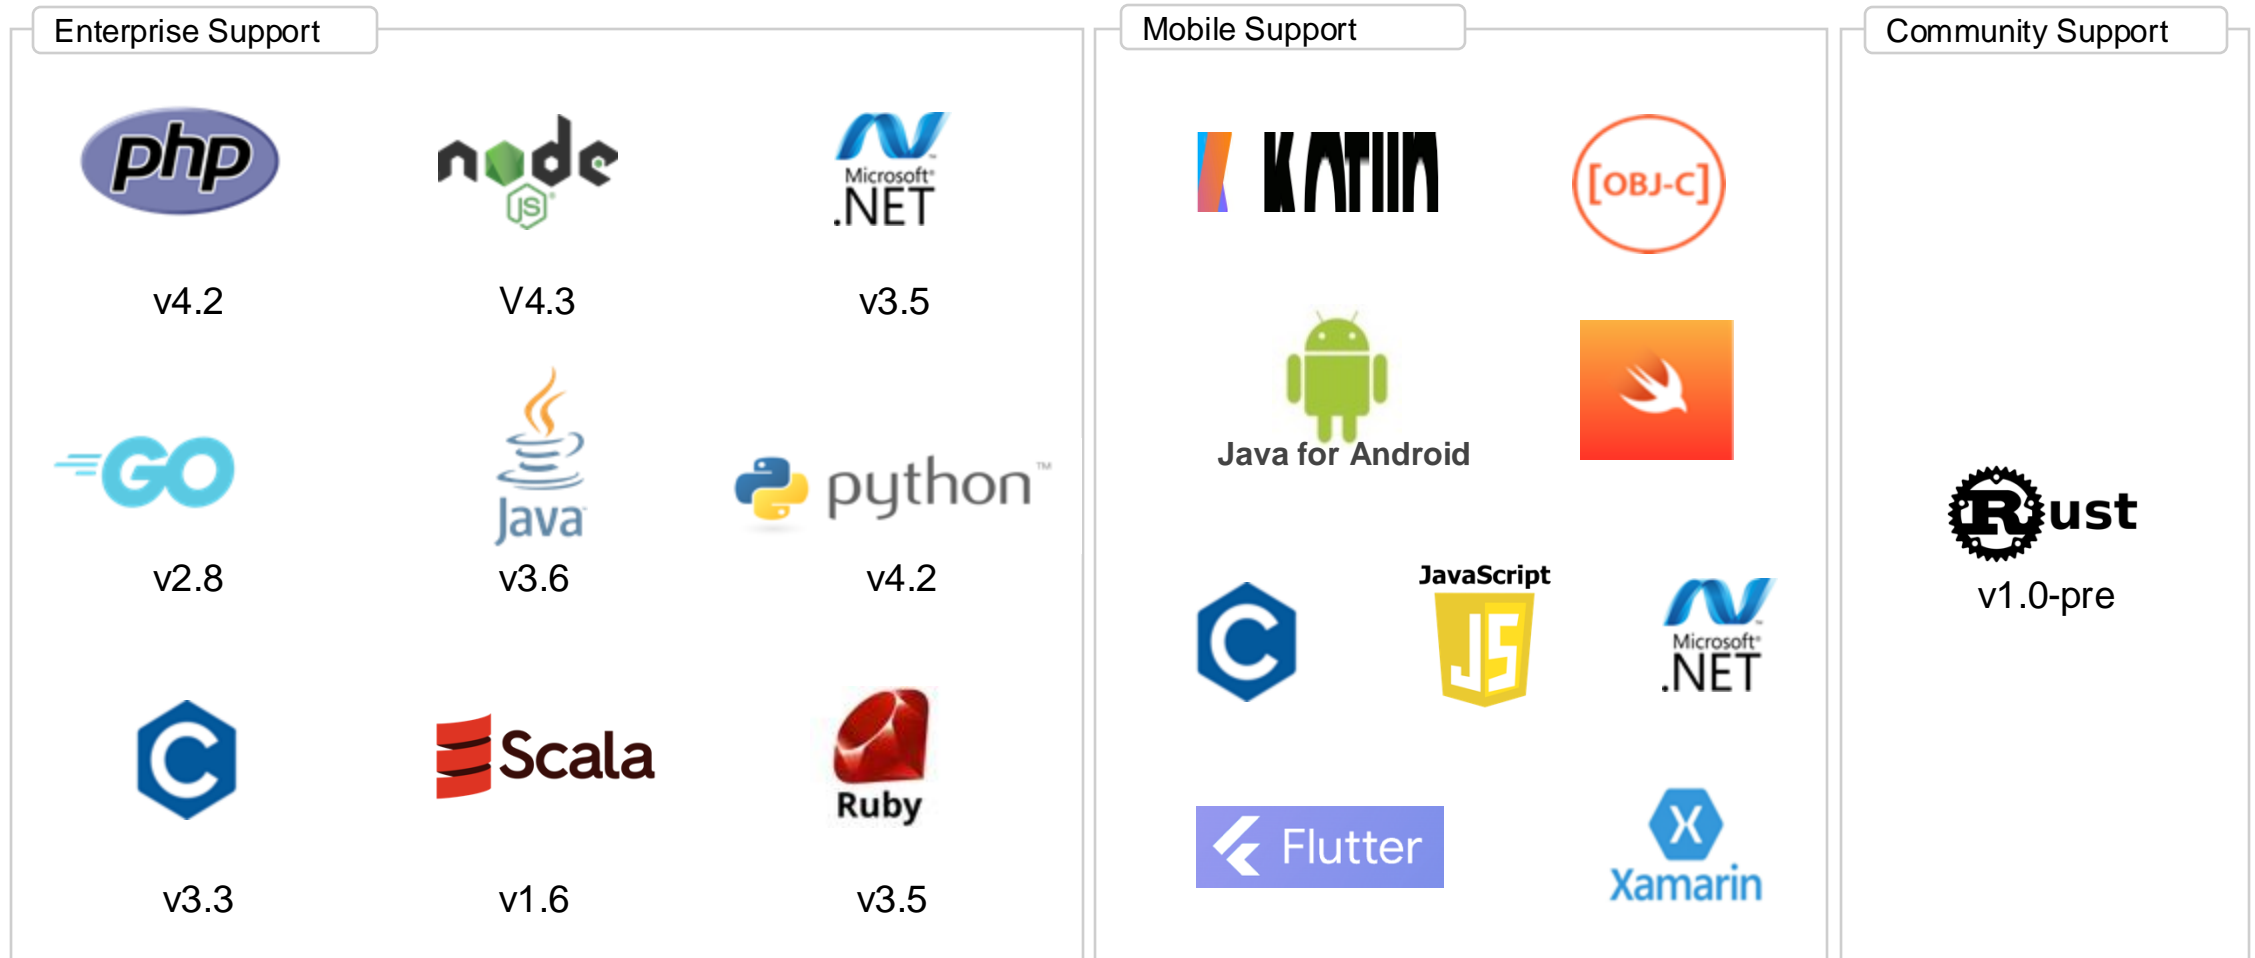


## 2-2. SDK 개요

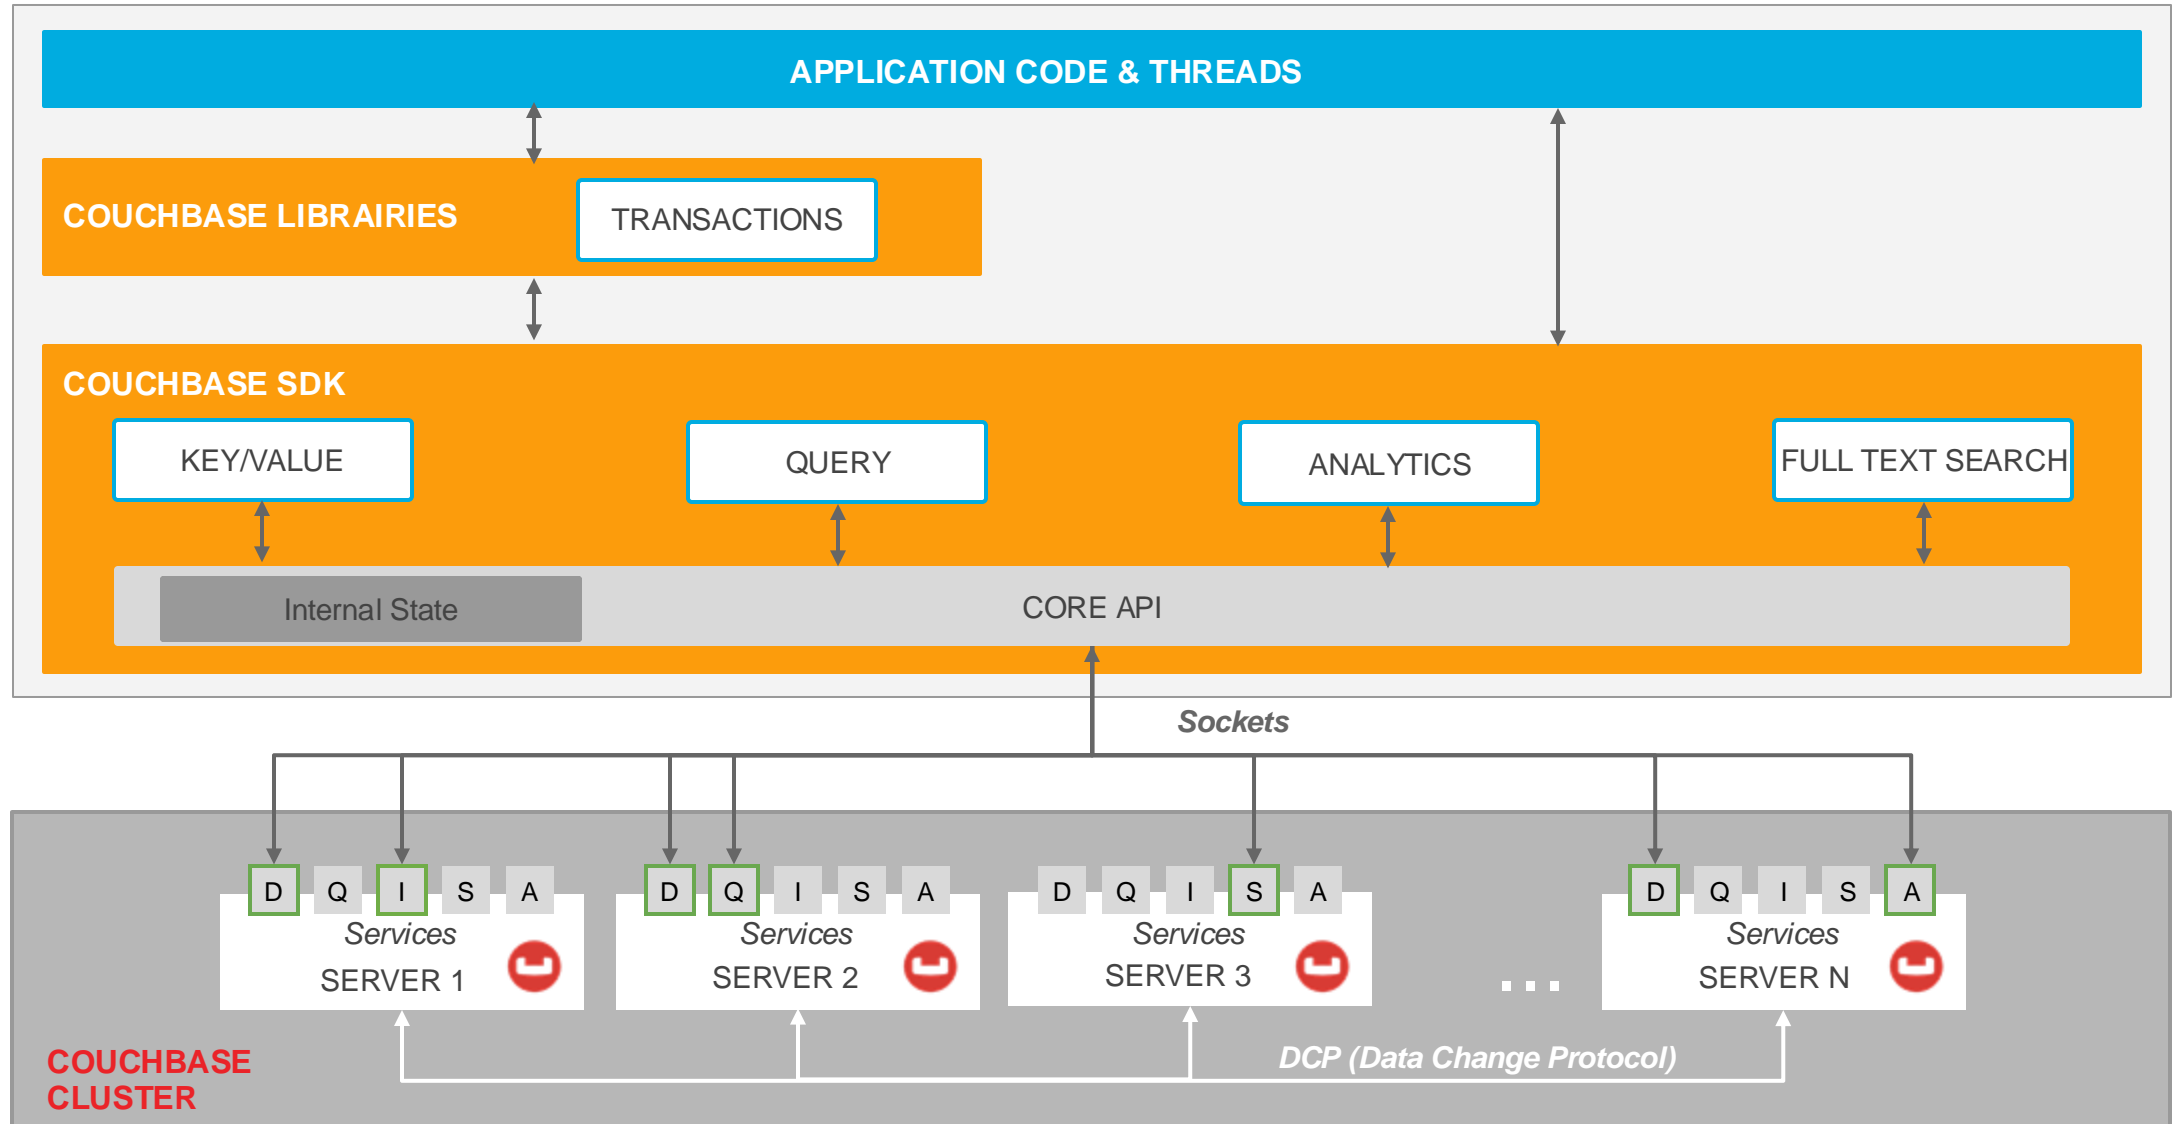


# SDK 개요

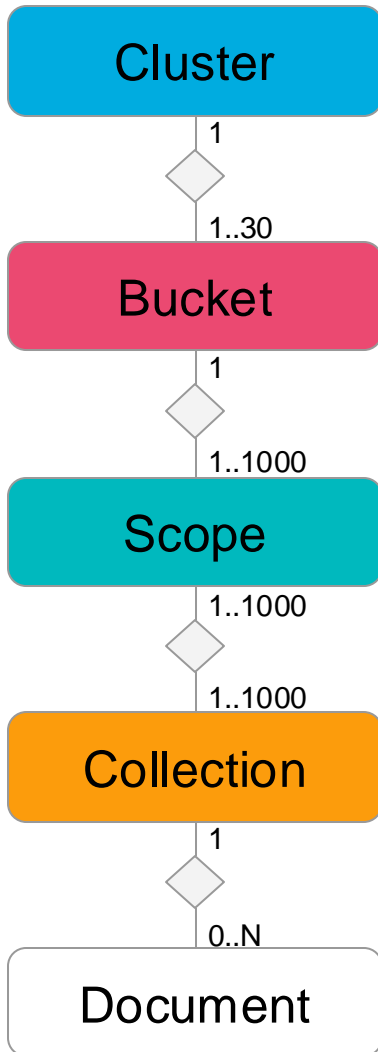
10 개 개발 언어에 대한 Couchbase 클러스터 접근을 위한 SDK(Software Development Kit) 제공



# SDK 아키텍처 : 통합 API 제공



# SDK 아키텍처 : 엔티티 계층구조



- Provides access to **cluster-level** operations(인스턴스)
- Manages bootstrapping, fault tolerance, rebalancing

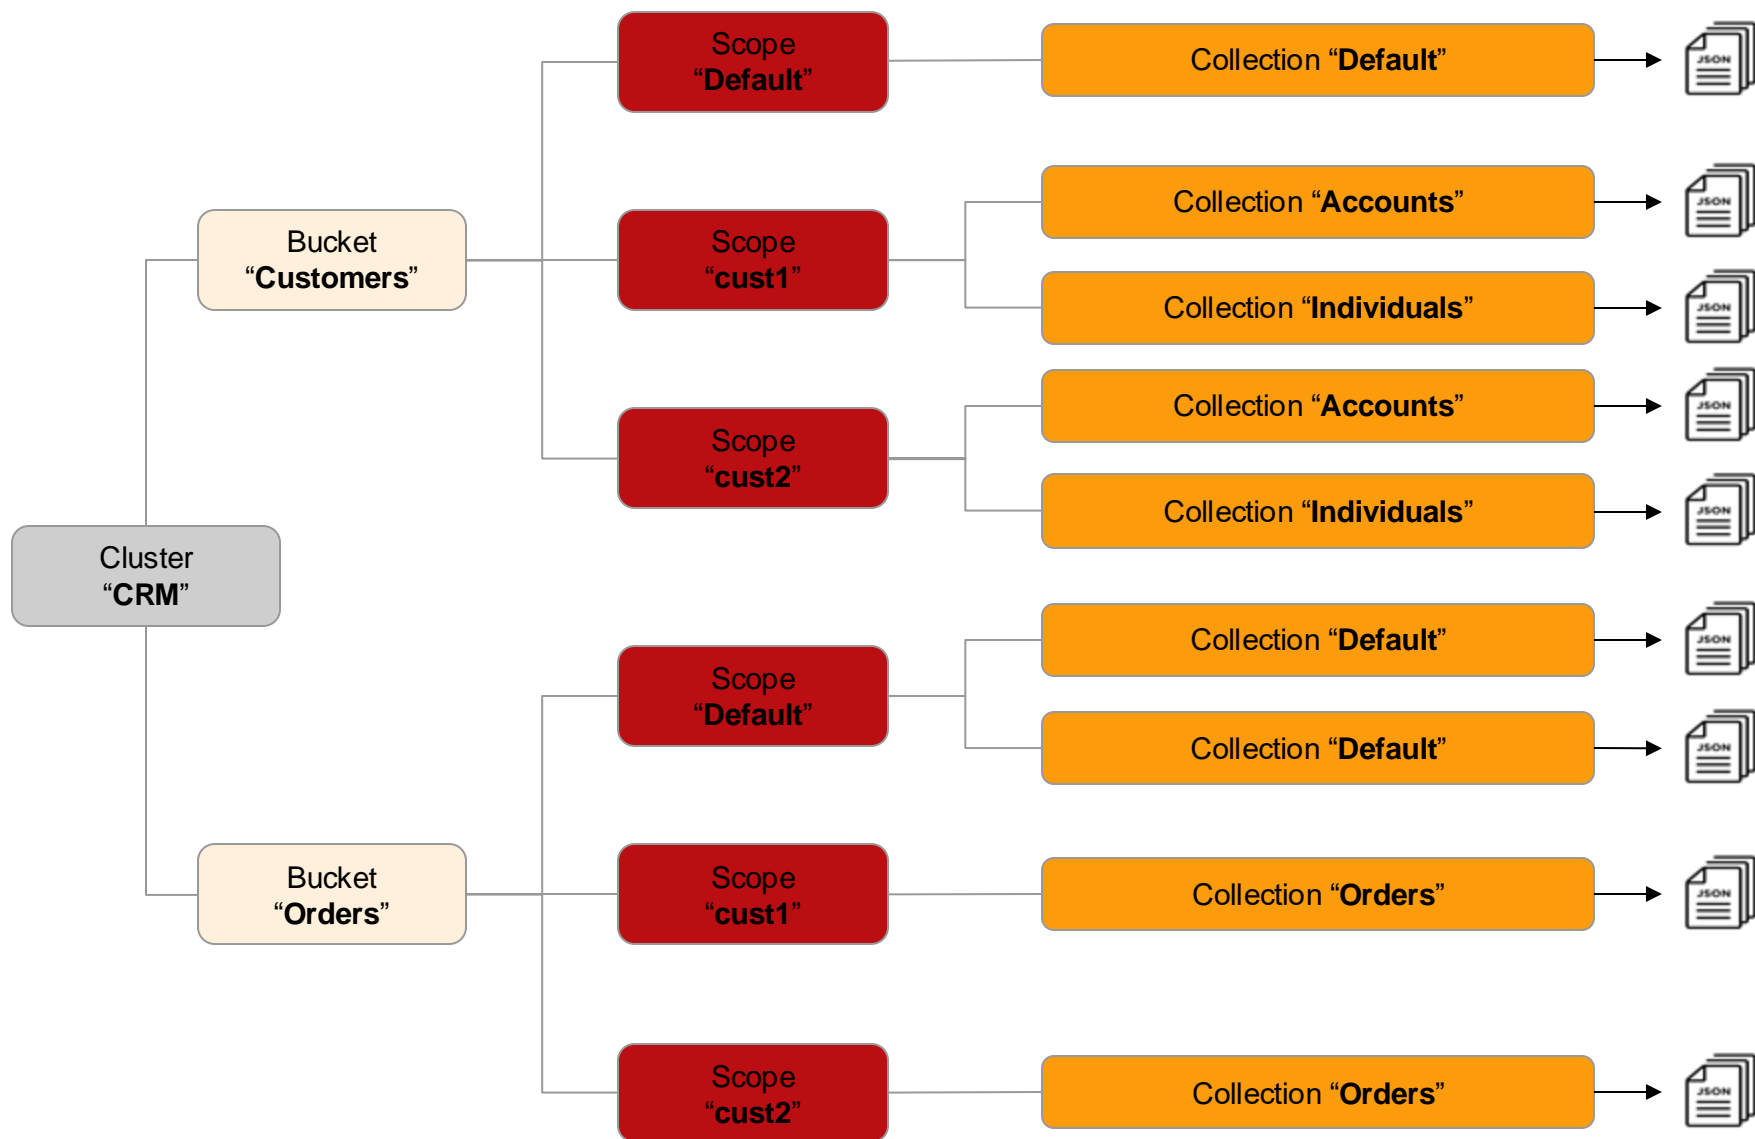
- Provides access to **bucket level** operations(데이터베이스)
- Manages persistence, replicas, TTL, compression mode, XDCR conflict resolution, ejection policy, compaction frequency

- Unique **keyspace of Collections**(스키마)
- Applications can be assigned per-scope access-rights

- Unique **keyspace of Documents**(테이블)
- Collections might be assigned to different scopes.

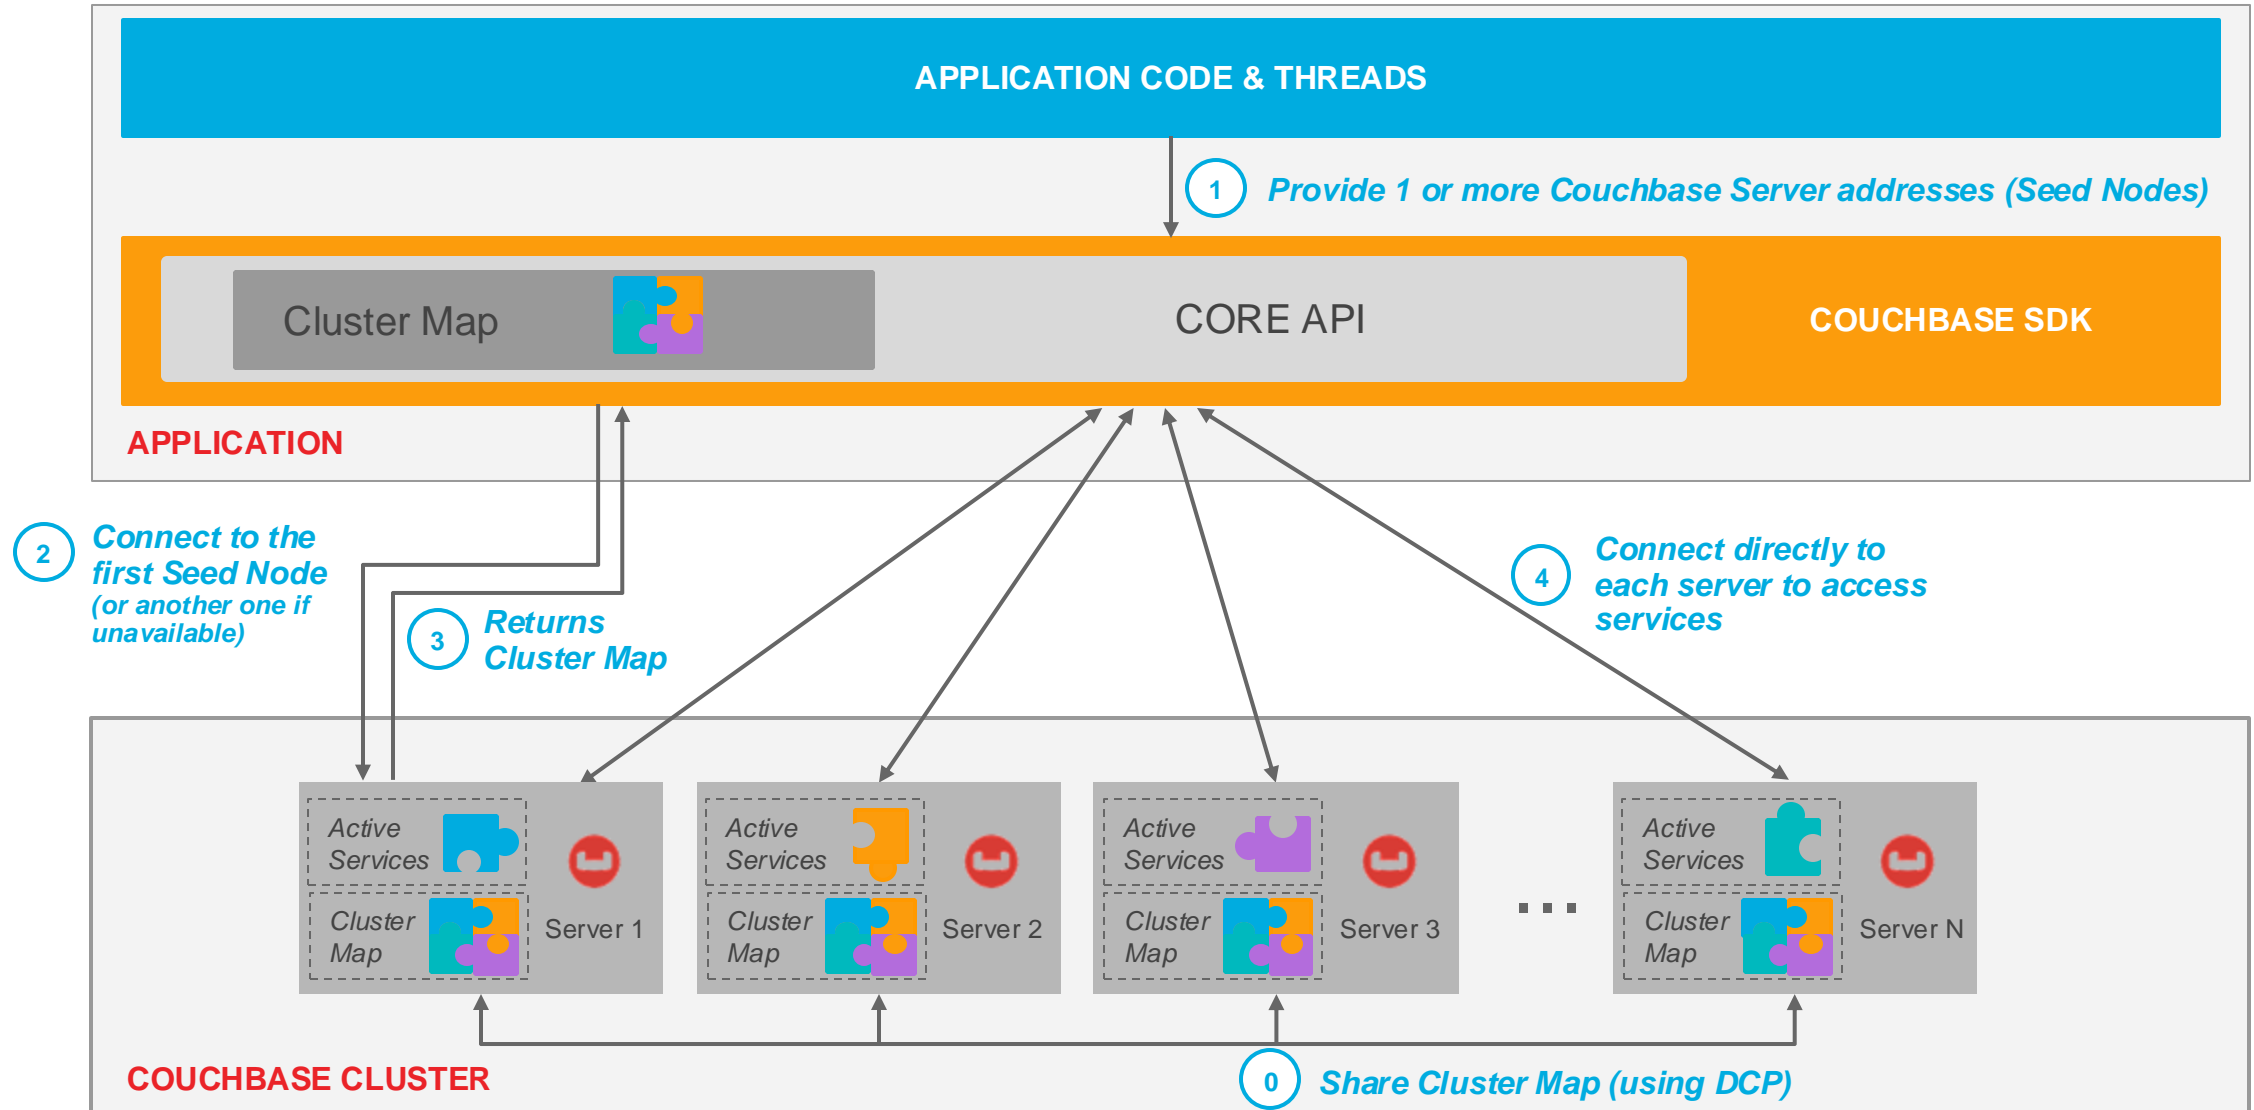
- JSON: numbers, strings, embedded documents and arrays(로우)
- Non JSON: UTF8, buffer, bytearray, serialized objects

# SDK 아키텍처 : 엔티티 계층구조





# SDK 아키텍처 : Connection 세부 단계(Bootstrapping)



# Sample Code : Synchronous

```
Cluster c = Cluster.connect("localhost", "username", "password");
Bucket b = c.bucket("travel-sample");
Collection col = b.defaultCollection();

// KV
MutationResult mr = col.upsert("doc1", JsonObject.create().put("name", "mike"));
GetResult d1 = col.get("doc1");
String name = d1.contentAsObject().getString("name");
System.out.println(name); // name == "mike"

// SQL++ query
final QueryResult qr = cluster.query("select * from `travel-sample` limit 10",
                                     queryOptions().metrics(true));
for (JsonObject row : qr.rowsAsObject()) {System.out.println("Row: " + row);}
System.out.println("Exec time:" + qr.metaData().metrics().get().executionTime());

// Search query
final SearchResult sr = cluster.searchQuery("airportIdx",
                                             SearchQuery.prefix("LAX"),
                                             searchOptions().fields("name"));
for (SearchRow row : sr.rows()){
    System.out.println("Score/ID: " + row.score("/") + row.id());
}

// Analytics query
AnalyticsResult ar = cluster.analyticsQuery(
    "select count(*) from airports where country = \"France\"");
for (JsonObject row : ar.rowsAsObject()) { System.out.println("Row: " + row);}
```

# Sample Code : Connections

```
static String endpoint = "couchbase://10.0.0.183,10.0.10.92";
static String username = "Administrator";
static String password = "P@ssw0rd";
static String bucketName = "travel-sample";
static String scopeName = "inventory";
static String collectionName = "airport";

ClusterEnvironment env = ClusterEnvironment.builder()
    .timeoutConfig(TimeoutConfig.kvTimeout(Duration.ofSeconds(15))
        .kvDurableTimeout(Duration.ofSeconds(15)))
    .thresholdLoggingTracerConfig(config)
    .build();

Cluster cluster = Cluster.connect(endpoint, ClusterOptions.clusterOptions(username, password)
    // Use the pre-configured profile below to avoid latency issues with your connection.
    .environment(env)
);

// get a bucket reference
Bucket bucket = cluster.bucket(bucketName);
bucket.waitForReady(Duration.ofSeconds(30));
Scope scope = bucket.scope(scopeName);
Collection collection = scope.collection(collectionName);
```

# 개발 지원 문서

The image displays three overlapping screenshots of the Couchbase documentation website, illustrating the structure and content of the developer resources.

**Left Screenshot:** Shows the Couchbase Documentation homepage with a navigation menu for SERVER, MOBILE, CAPELLA, and CLOUD-NATIVE. A sidebar lists various SDKs and connectors with version numbers:

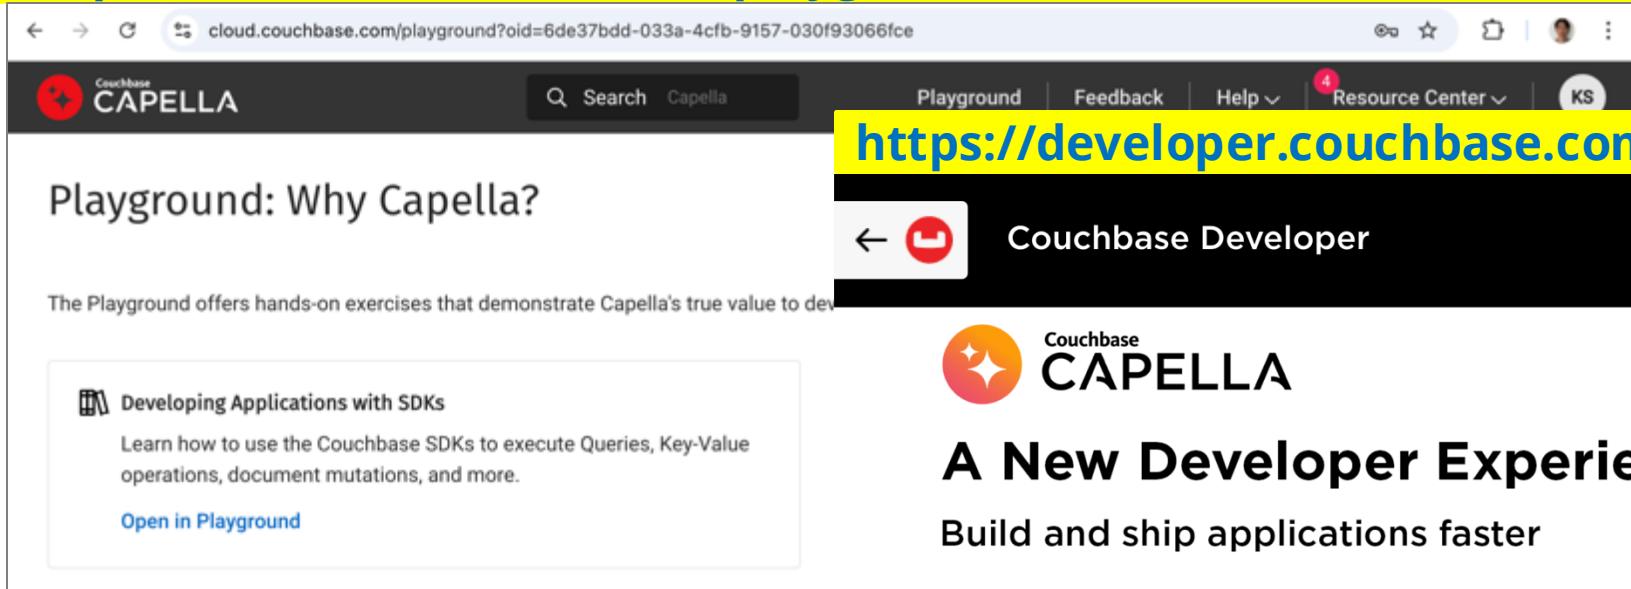
- .NET SDK (3.4)
- C SDK
- Go SDK (2.7)
- Java SDK (3.5)
- Kotlin SDK (1.2)
- Node.js SDK (4.2)
- PHP SDK (4.1)
- Python SDK (4.1)
- Ruby SDK (3.4)
- Scala SDK (1.5)
- C++ Transactions
- Elasticsearch Connector
- Kafka Connector (4.2)
- Spark Connector (3.3)
- Tableau Connector
- Power BI Connector

**Middle Screenshot:** Shows the "Data Operations" section under the "SDKS" category. It includes a "Getting Started" guide, a "Data Operations" overview, and a "Documents" section explaining document structure and keys.

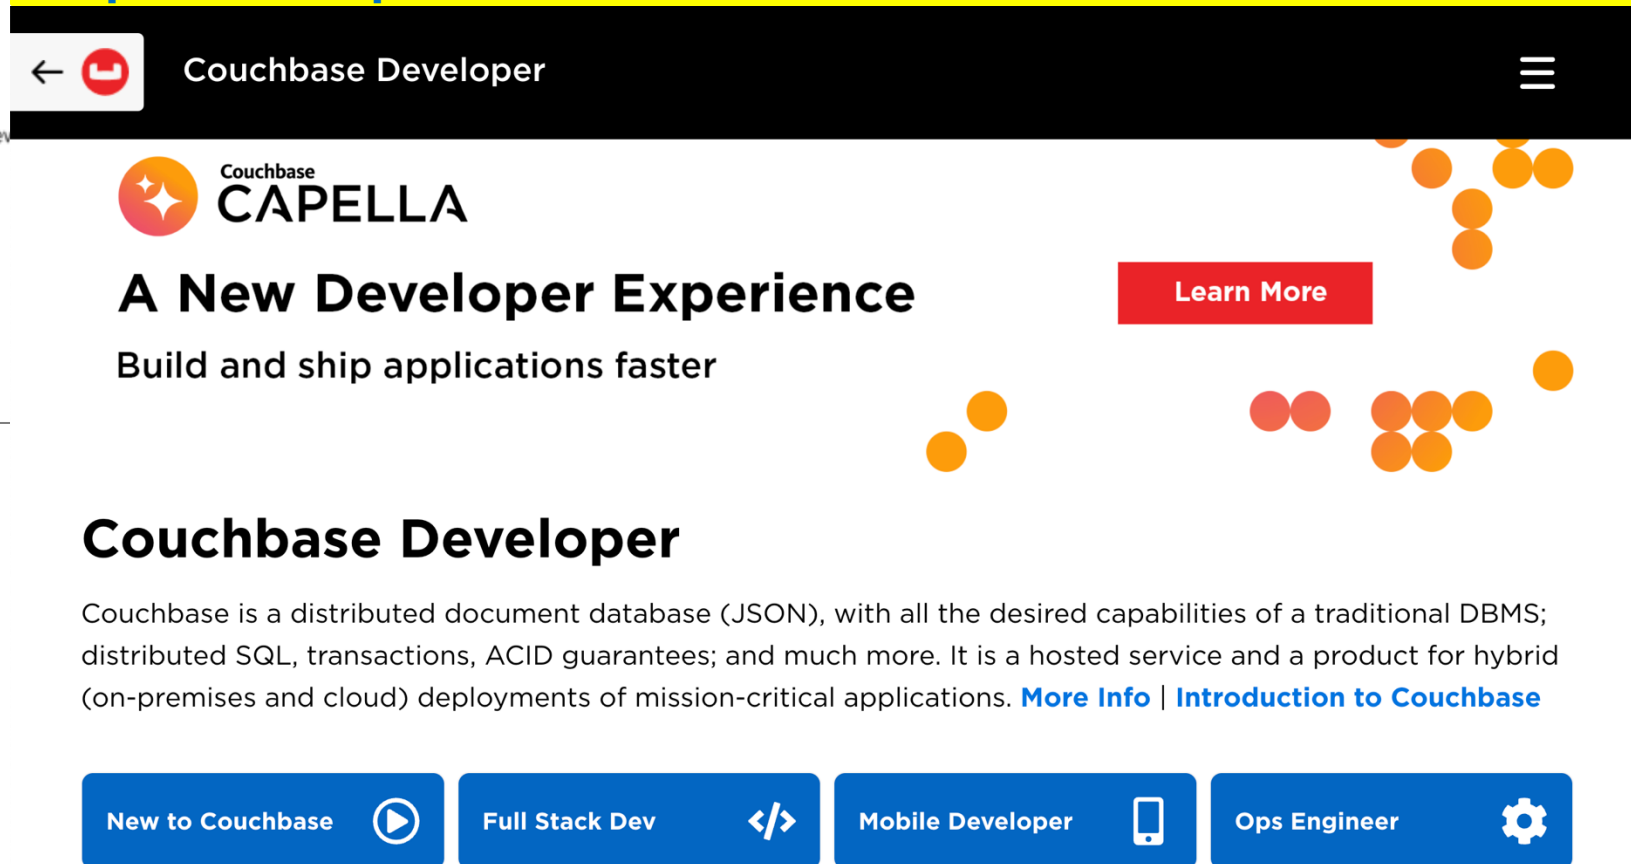
**Right Screenshot:** Shows the "Sample Application" section, which provides a "Quick Start" guide. It includes instructions on how to fetch the Couchbase .NET SDK travel-sample Application REST Backend from GitHub and how to run the application using Docker.

# 개발자 포털/플레이그라운드

<https://cloud.couchbase.com/playground>



<https://developer.couchbase.com/tutorials>

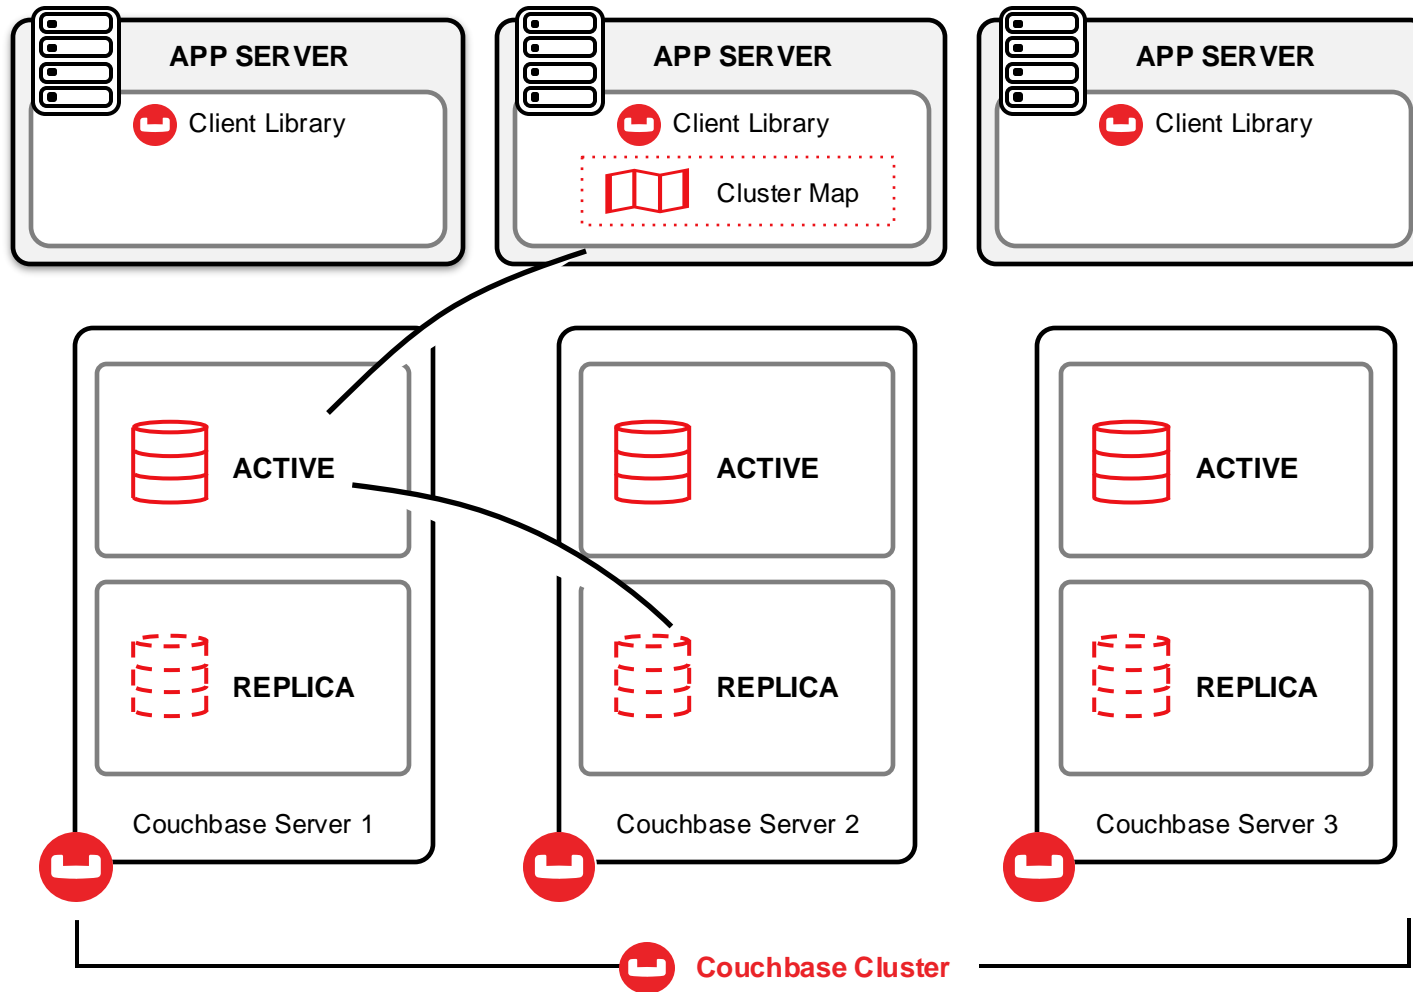


## 2-3. SDK 실습 (Python, Java)



The screenshot shows the Couchbase Documentation website. The top navigation bar includes the Couchbase logo, 'Documentation', a search bar, 'Downloads', and a 'Try Free' button. Below this is a secondary navigation bar with links for SERVER, MOBILE, CAPELLA, CLOUD-NATIVE, COUCHBASE SDKS (highlighted in red), and COLUMNAR SDKS. The left sidebar shows the 'Python SDK' section with a version dropdown set to '4.3'. Under 'Getting Started', 'Start Using the Python SDK' is selected. The main content area is titled 'Start Using the Python SDK' with a 'TUTORIAL' badge. It includes a sub-header 'Python SDK / Getting Started / Start Using the Python SDK', an 'Edit on GitHub' link, and a paragraph: 'Get up and running quickly, installing the Couchbase Python SDK, and running our Hello World example.' Below this, it states: 'The Couchbase Python SDK allows Python applications to access a Couchbase cluster. It offers a traditional synchronous API as well as integration with *twisted* and *asyncio*.' A section 'In this guide, you will learn:' lists three bullet points: 'How to connect to Couchbase Capella or Couchbase Server.', 'How to add and retrieve Documents.', and 'How to lookup documents with the SQL++ (formerly N1QL) query language.' The next section is 'Hello Couchbase', which says: 'We will go through the code sample step by step, but for those in a hurry to see it, here it is:'. The right sidebar contains a list of links: 'Hello Couchbase', 'Quick Installation', 'Prerequisites', 'Step-by-Step' (with sub-links: '- Connect', '- Add and Retrieve Documents', '- SQL++ Lookup', '- Execute!'), 'Next Steps' (with sub-links: '- Additional Resources', '- Troubleshooting'), 'Is this page helpful?' (with 'Yes' and 'No' buttons), and 'Leave Additional Feedback?'. The footer of the page contains the text: 'Confidential and Proprietary. Do not distribute without Couchbase consent. © Couchbase 2024. All rights reserved.'

# A Fully Distributed Database



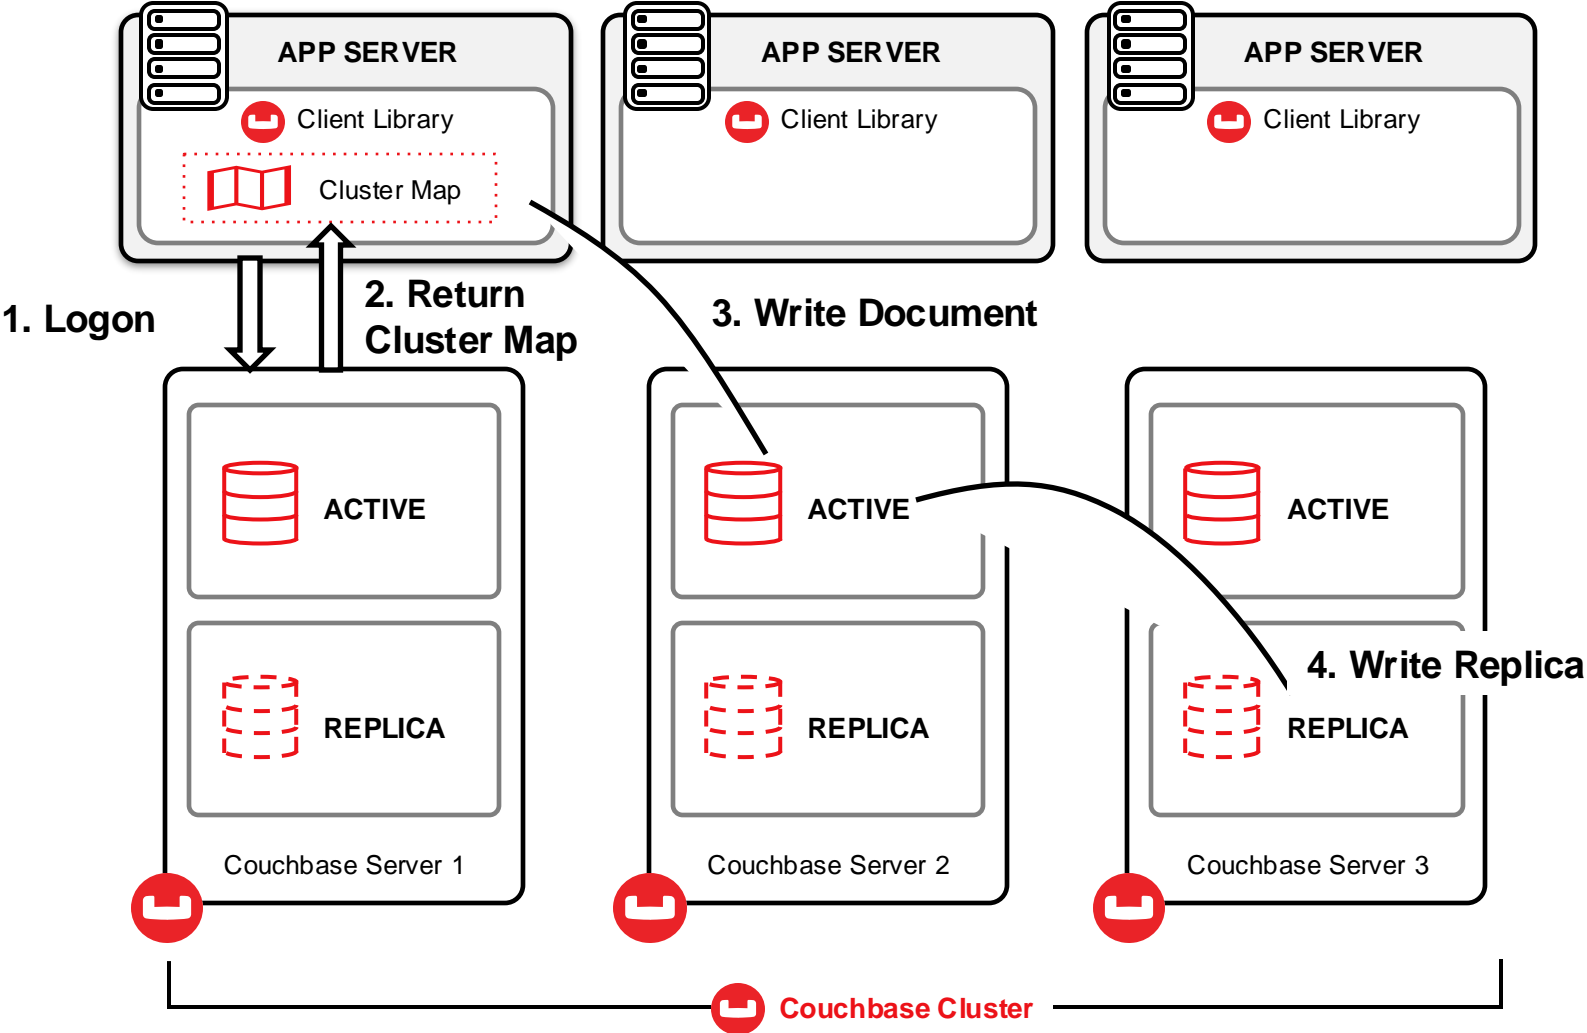
**No Master-Slave**  
Single node type  
Direct client communication

**Automatic sharding**  
Maximum cluster efficiency

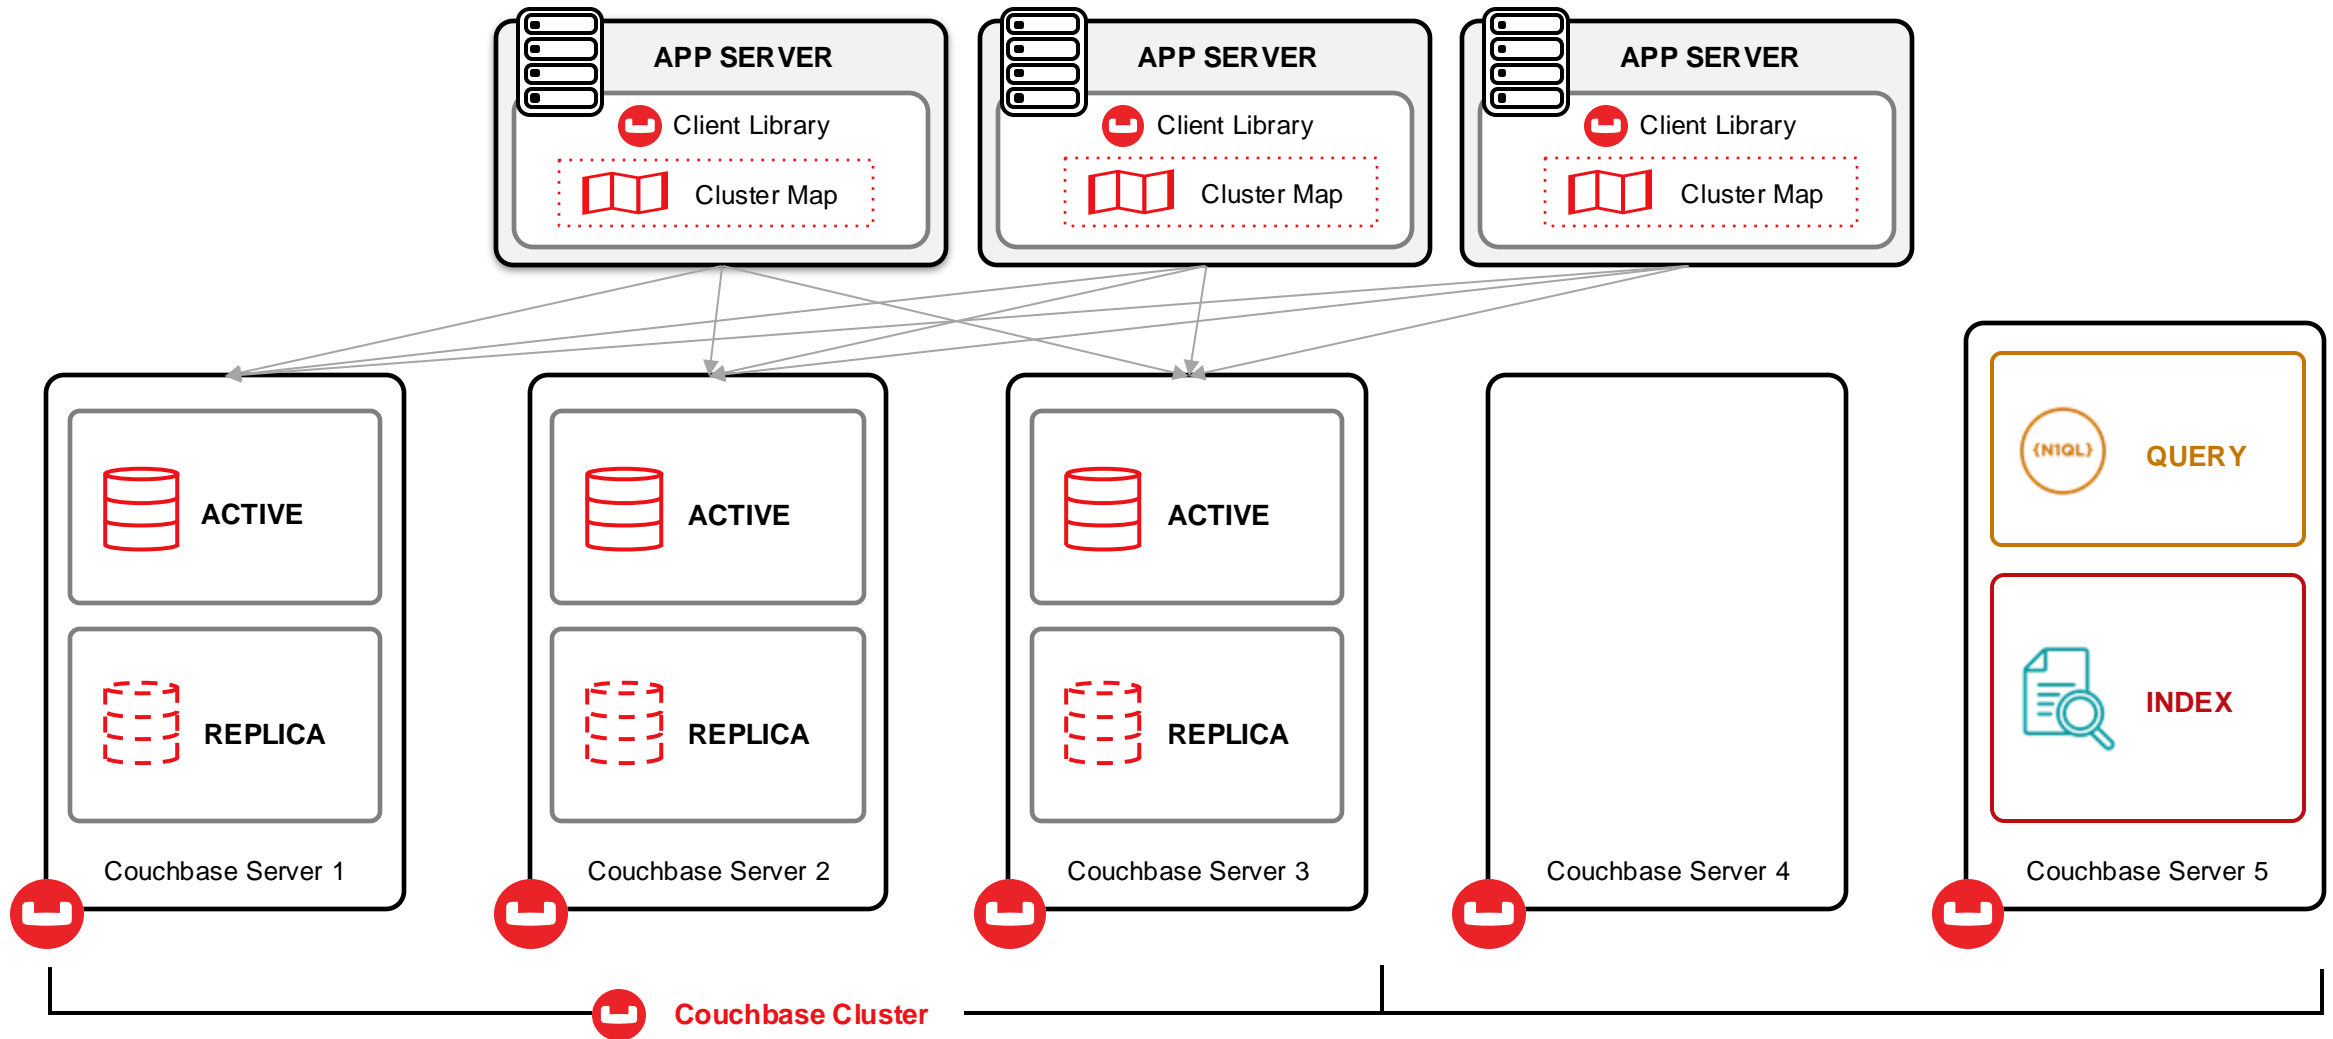
**Automatic failover**  
Continuous replication  
for no disruption



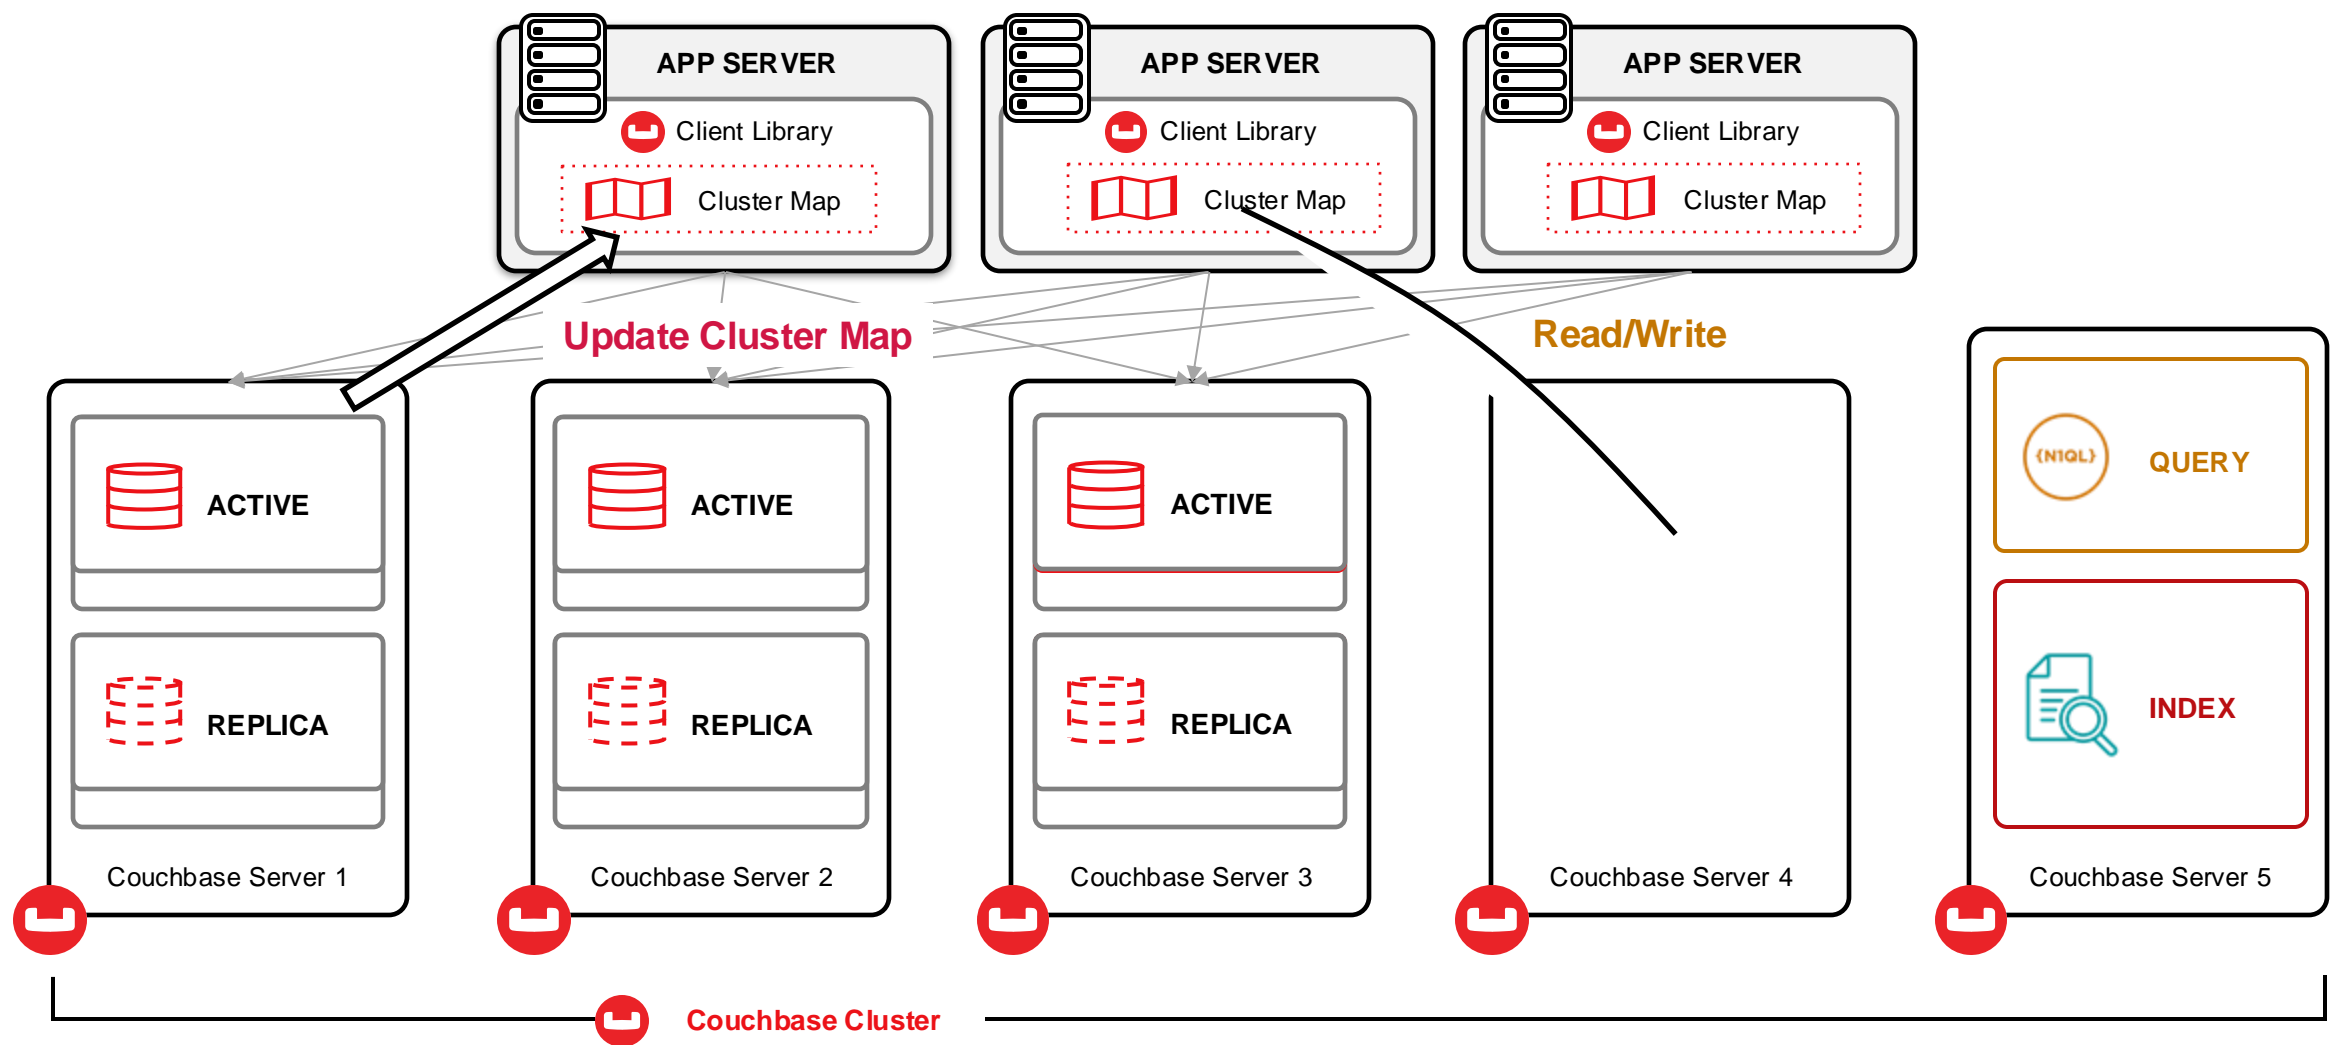
# Basics Operations



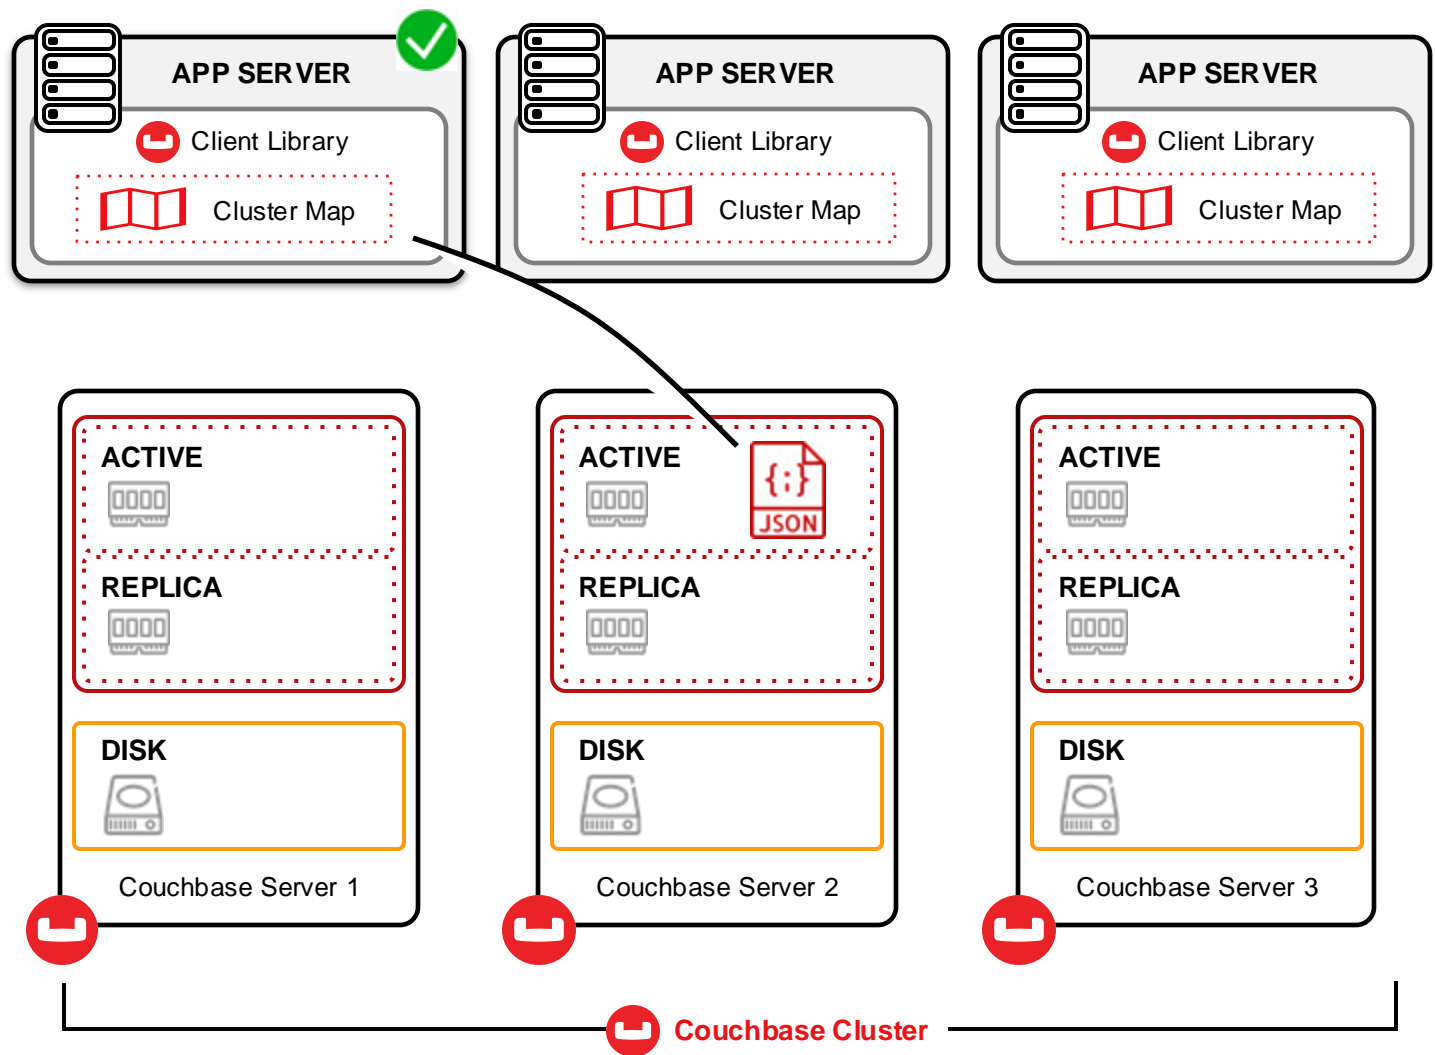
# Scalability – Horizontal & Vertical



# Scalability – Horizontal & Vertical



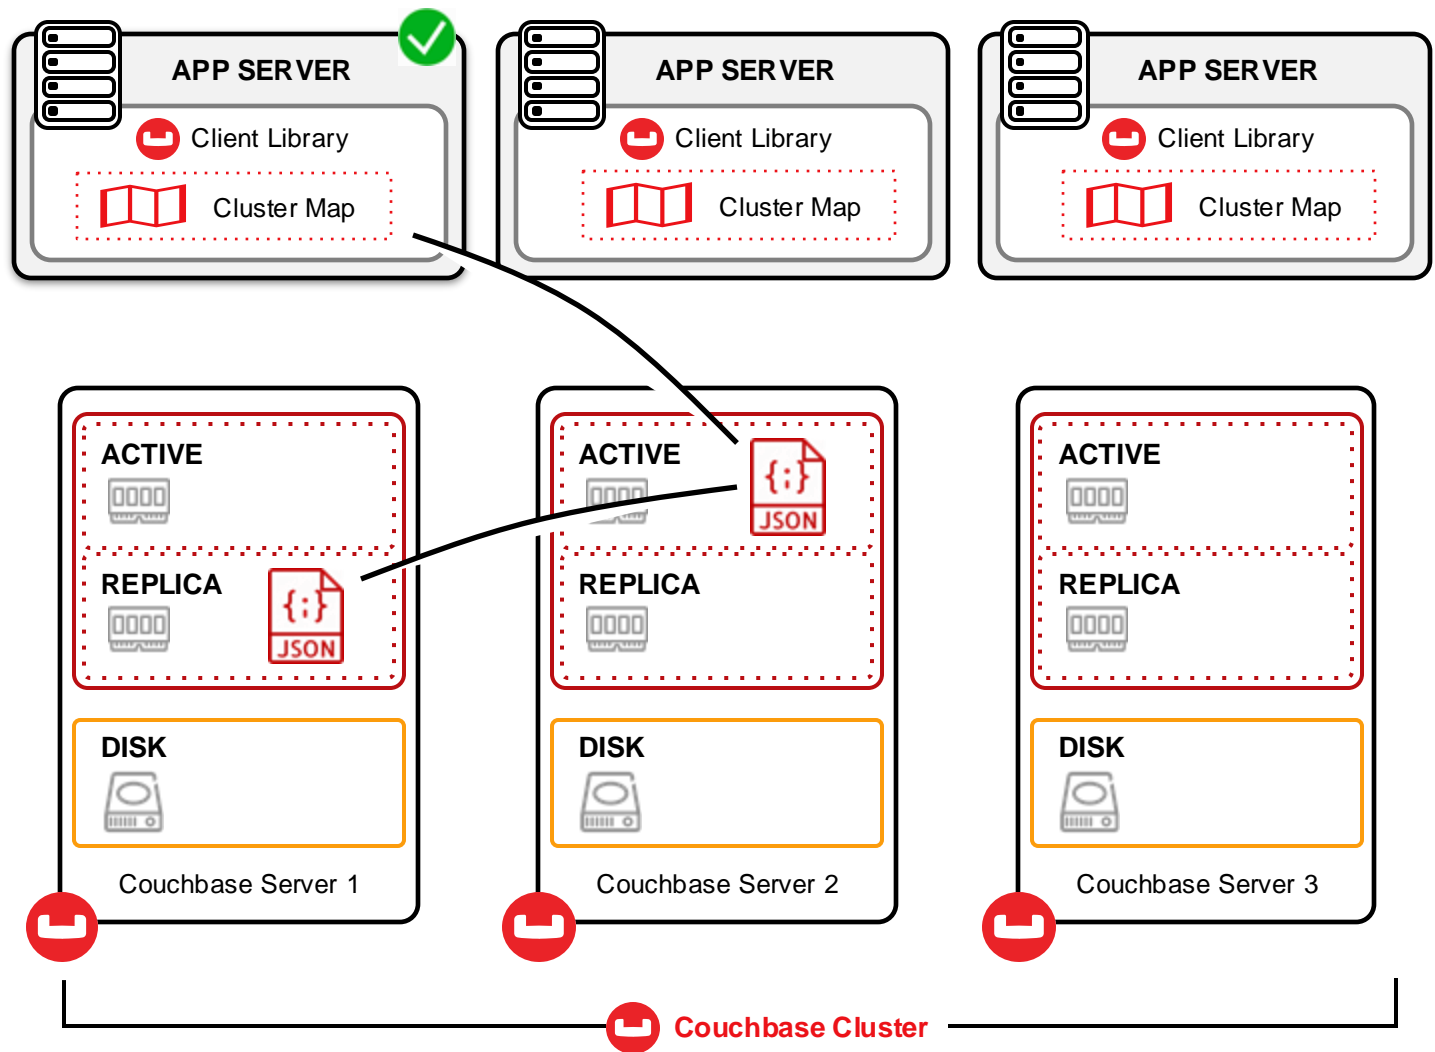
# Durability



replica = 1

Level of Durability
None
Majority
MajorityAndPersistToActive
PersistToMajority

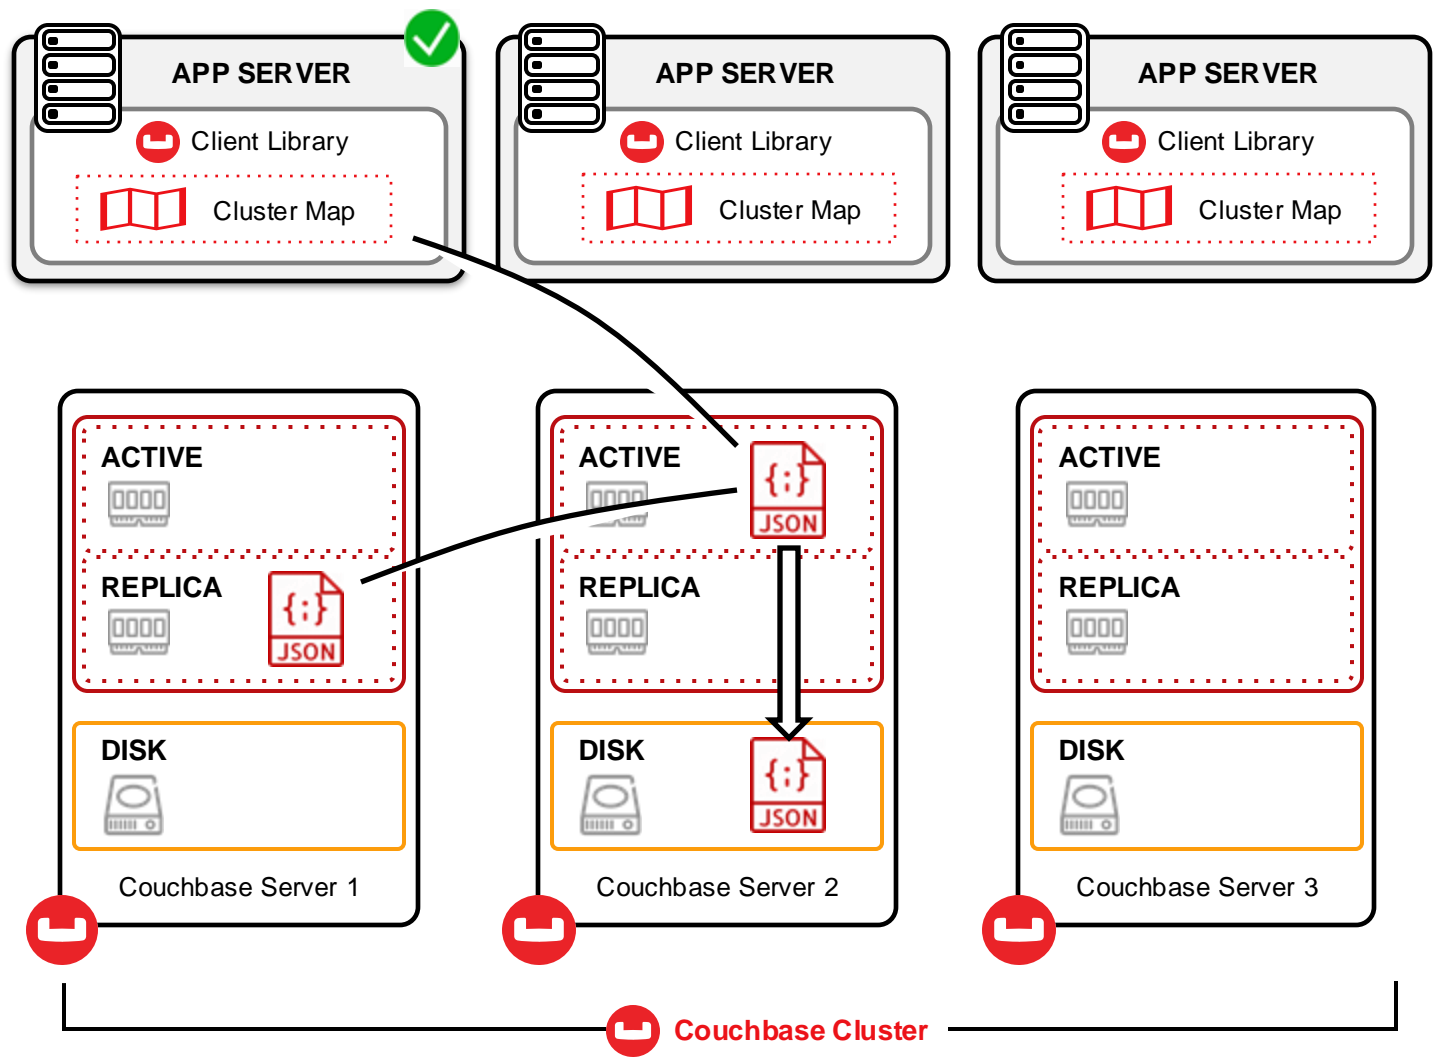
# Durability



replica = 1

Level of Durability
None
<b>Majority</b>
MajorityAndPersistToActive
PersistToMajority

# Durability



replica = 1

## Level of Durability

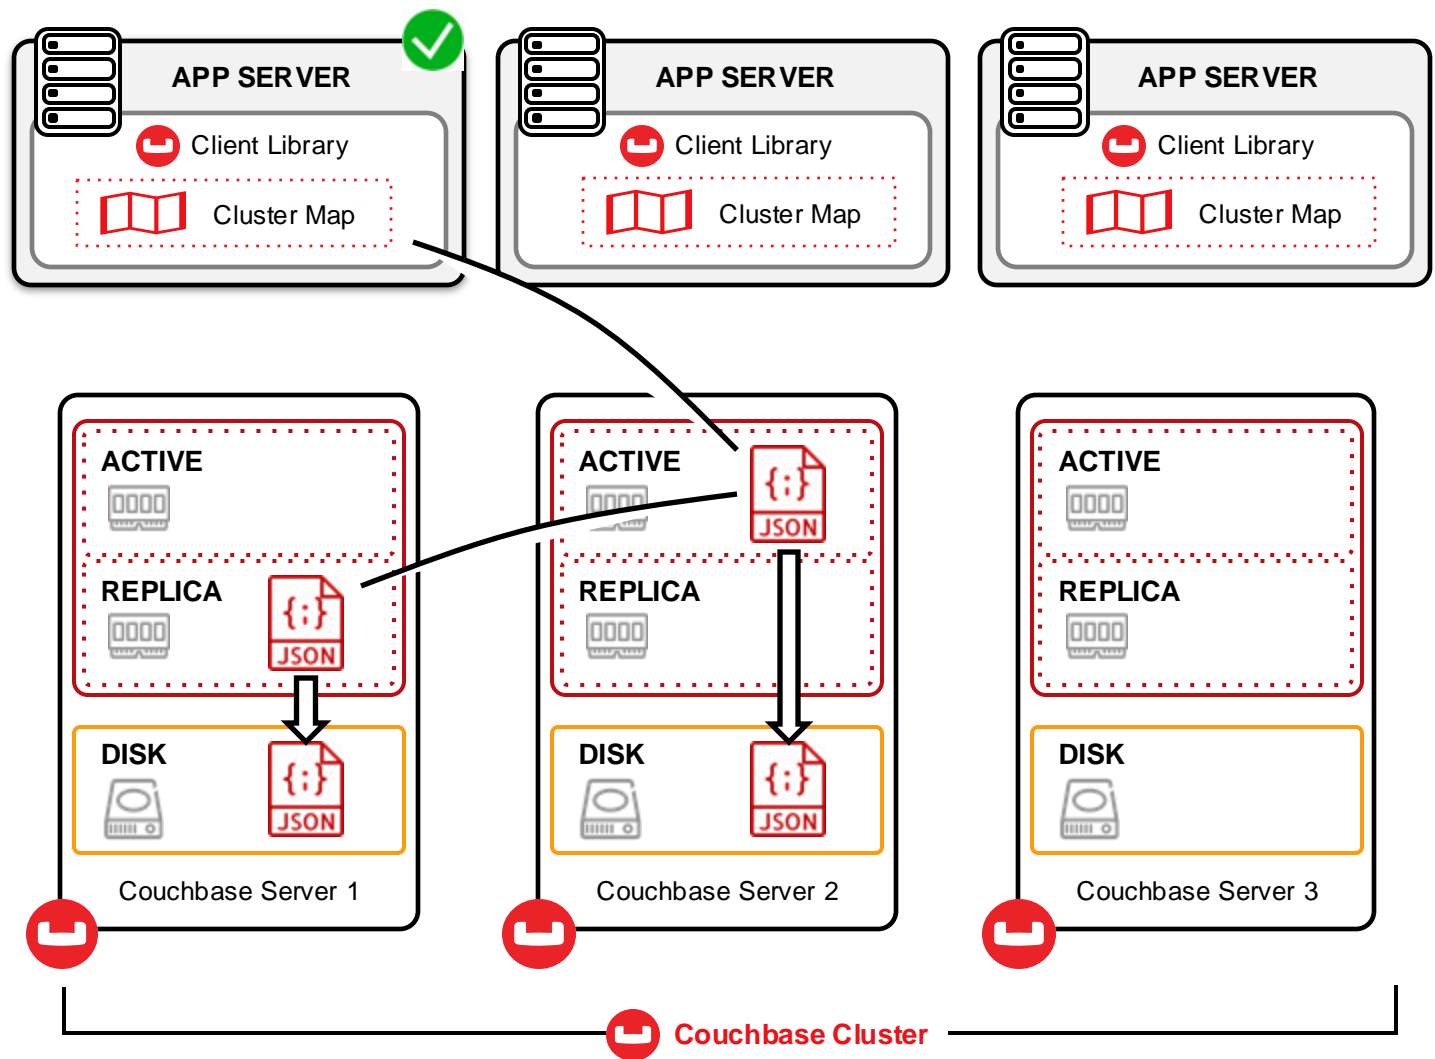
None

Majority

**MajorityAndPersistToActive**

PersistToMajority

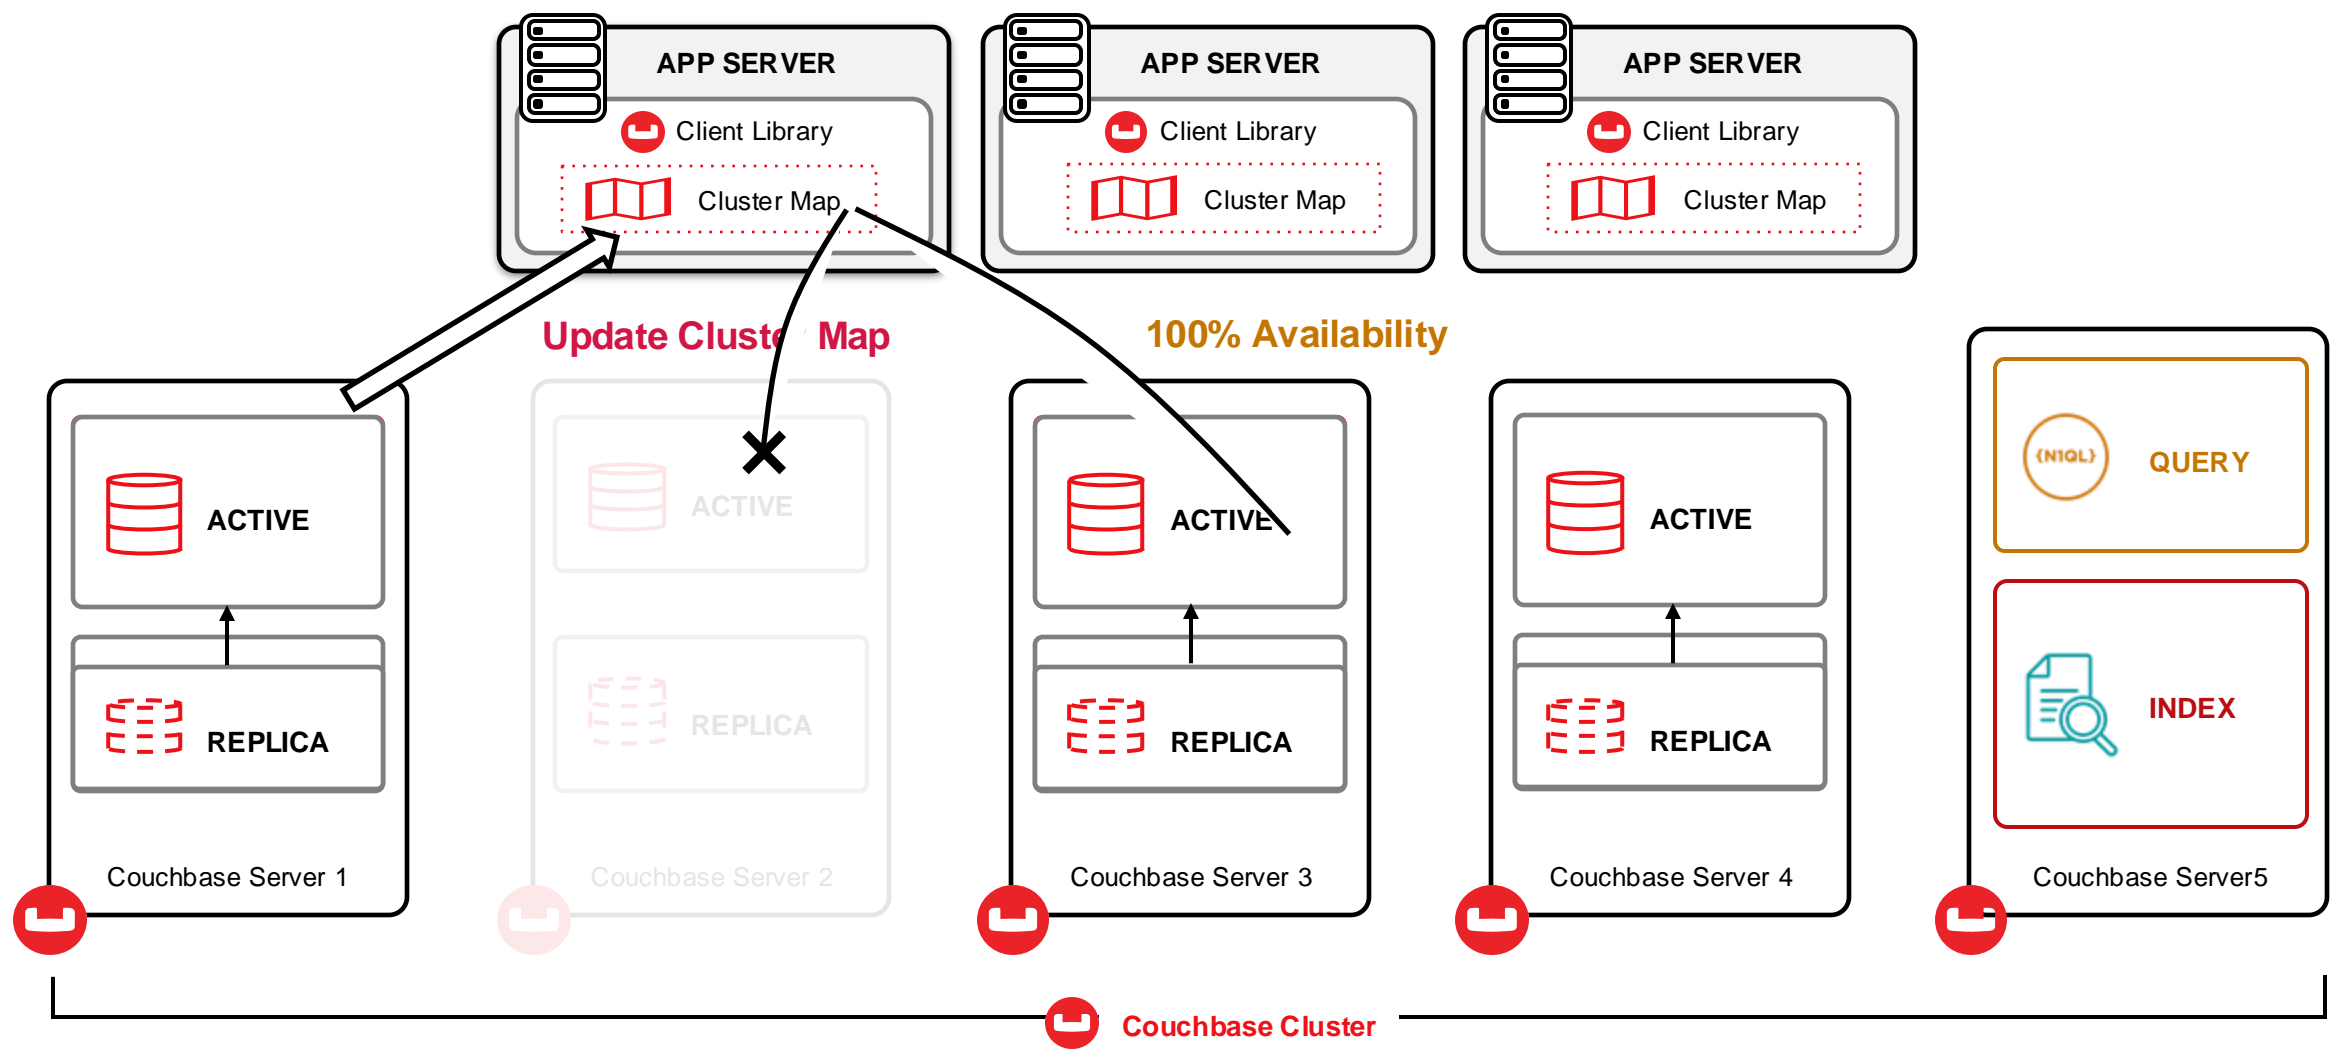
# Durability



replica = 1

Level of Durability
None
Majority
MajorityAndPersistToActive
<b>PersistToMajority</b>

# Node Failover





# SQL Sample from Capella Playground

```
SELECT route.airline, route.schedule, route.sourceairport
FROM route
WHERE route.sourceairport = "JFK"
ORDER BY route.airline LIMIT 5;
```

```
SELECT DISTINCT airline.name, airline.callsign, route.schedule,
route.sourceairport, route.destinationairport
FROM route
INNER JOIN airline
ON route.airlineid = META(airline).id
WHERE route.sourceairport = "JFK"
AND route.destinationairport = "SFO"
AND airline.callsign = "UNITED";
```

```
SELECT h.name, h.address
FROM hotel h
WHERE h.city = 'San Francisco'
AND ANY r IN h.reviews
SATISFIES r.ratings.Overall > 3 END
```

```
SELECT h.city, AVG(r.ratings.Overall) AS hotel_rating
FROM hotel AS h
UNNEST h.reviews AS r
WHERE h.country = "United Kingdom"
GROUP BY h.city
HAVING AVG(r.ratings.Overall) > 4.9
```

```
SELECT h.city, h.name,
AVG(r.ratings.Overall) AS avg_rating,
RANK() OVER (PARTITION BY h.city ORDER BY AVG(r.ratings.Overall) DESC) as city_rank
FROM hotel as h
UNNEST h.reviews r
WHERE h.country = "France"
GROUP BY h.city, h.name
```

# SQL Sample from Capella Playground

```
WITH AirlineDestinations AS (  
  SELECT r.destinationairport,  
  ARRAY_COUNT(r.schedule) AS flights  
  FROM `travel-sample`.inventory.route AS r  
  WHERE r.airline = "UA"  
  GROUP BY r.destinationairport, r.schedule  
)  
SELECT ad.destinationairport, SUM(ad.flights) AS total_flights  
FROM AirlineDestinations AS ad  
GROUP BY ad.destinationairport  
ORDER BY ad.destinationairport;
```

```
select a.name, count(1) as numRoutes  
from route r  
join airline a on r.airlineid = meta(a).id  
group by a.name  
having count(1) > 2000  
order by count(1) desc;
```

```
SELECT h.city,  
  ARRAY_AGG(h.name)  
  FILTER (WHERE ANY AND EVERY r IN h.reviews  
    SATISFIES r.ratings.Overall >= 3 END) AS good_hotels,  
  ARRAY_AGG(h.name)  
  FILTER (WHERE ANY r IN h.reviews  
    SATISFIES r.ratings.Overall < 3 END) AS other_hotels  
FROM hotel as h  
GROUP BY h.city;
```

```
select meta().id as _id, * from route
```

# Appendix. SQL Samples



# SQL++ Queries design

Modelling data using Collections simplifies queries and indexes

```
CREATE BUCKET CRM;  
set query_context = "default:CRM._default"
```

Set query context. Here \_default scope of bucket CRM.

```
CREATE COLLECTION Orders;  
CREATE COLLECTION Individuals;
```

Create collections in that context

```
INSERT INTO Individuals (key,value)  
VALUES ("Ind:110", { "Id": "110","Name": "Mary Joe",  
                    "Email": "mj@email.com" });  
  
INSERT INTO Orders (key,value)  
VALUES ("Ord:123", { "Id": "123", "Indid": "110",  
                    "Items": [ "Qty": "1","Name": "Shoes XX"}] });
```

Insert Json docs into collections.  
There is no type field.

```
CREATE INDEX Orders_Indid ON Orders(Indid);  
CREATE INDEX Individuals_Id ON Individuals(Id);
```

Create indexes into a collection.  
There is no WHERE type='Orders'.

```
SELECT i.Name, o.Items  
FROM Orders o  
JOIN Individuals i ON i.Id = o.Indid;
```

Query with JOIN between collections

# SQL-like queries | MERGE

***Finds all routes for airline BA whose source airport is in France.***

- *If any flights are using equipment 319, they are updated to use equipment 797.*
- *If any flights are using equipment 757, they are deleted.*

Collection: 'route'

```
{
  "id": "14452",
  "type": "route",
  "airline": "BA",
  "airlineid": "airline_1355",
  "sourceairport": "BOD",
  "destinationairport": "LGW",
  "equipment": "734 319",
  "stops": "0",
  "distance": 99.89861063028253
  "schedule": [
    {"day": 0, "utc": "19:41:00", "flight": "BA977"},
    {"day": 0, "utc": "14:03:00", "flight": "BA938"},
    {"day": 1, "utc": "08:40:00", "flight": "BA394"},
    {"day": 1, "utc": "06:54:00", "flight": "BA427"},
    {"day": 1, "utc": "06:00:00", "flight": "BA916"}
  ]
}
```

Collection: 'airport'

```
{
  "id": 1264,
  "type": "airport",
  "airportname": "Merignac",
  "city": "Bordeaux",
  "country": "France",
  "faa": "BOD",
  "icao": "LFBD",
  "tz": "Europe/Paris",
  "geo": {
    "lat": 44.828335,
    "lon": -0.715556,
    "alt": 162
  }
}
```

***Before updating or deleting*** some routes,  
***you first need to JOIN*** route and airport because the  
*country of the source airport is in the airport collection*

# SQL-like queries | MERGE

```
MERGE INTO `travel-sample`.inventory.route
USING `travel-sample`.inventory.airport
ON route.sourceairport = airport.faa
WHEN MATCHED THEN
  UPDATE
    SET route.old_equipment = route.equipment,
        route.equipment = "797"
    WHERE airport.country = "France" AND route.airline = "BA" AND CONTAINS(route.equipment, "319")
WHEN MATCHED THEN
  DELETE
    WHERE airport.country = "France" AND route.airline = "BA" AND CONTAINS(route.equipment, "757")
RETURNING route.old_equipment, route.equipment, airport.faa;
```

Collection to modify

Collection to join with

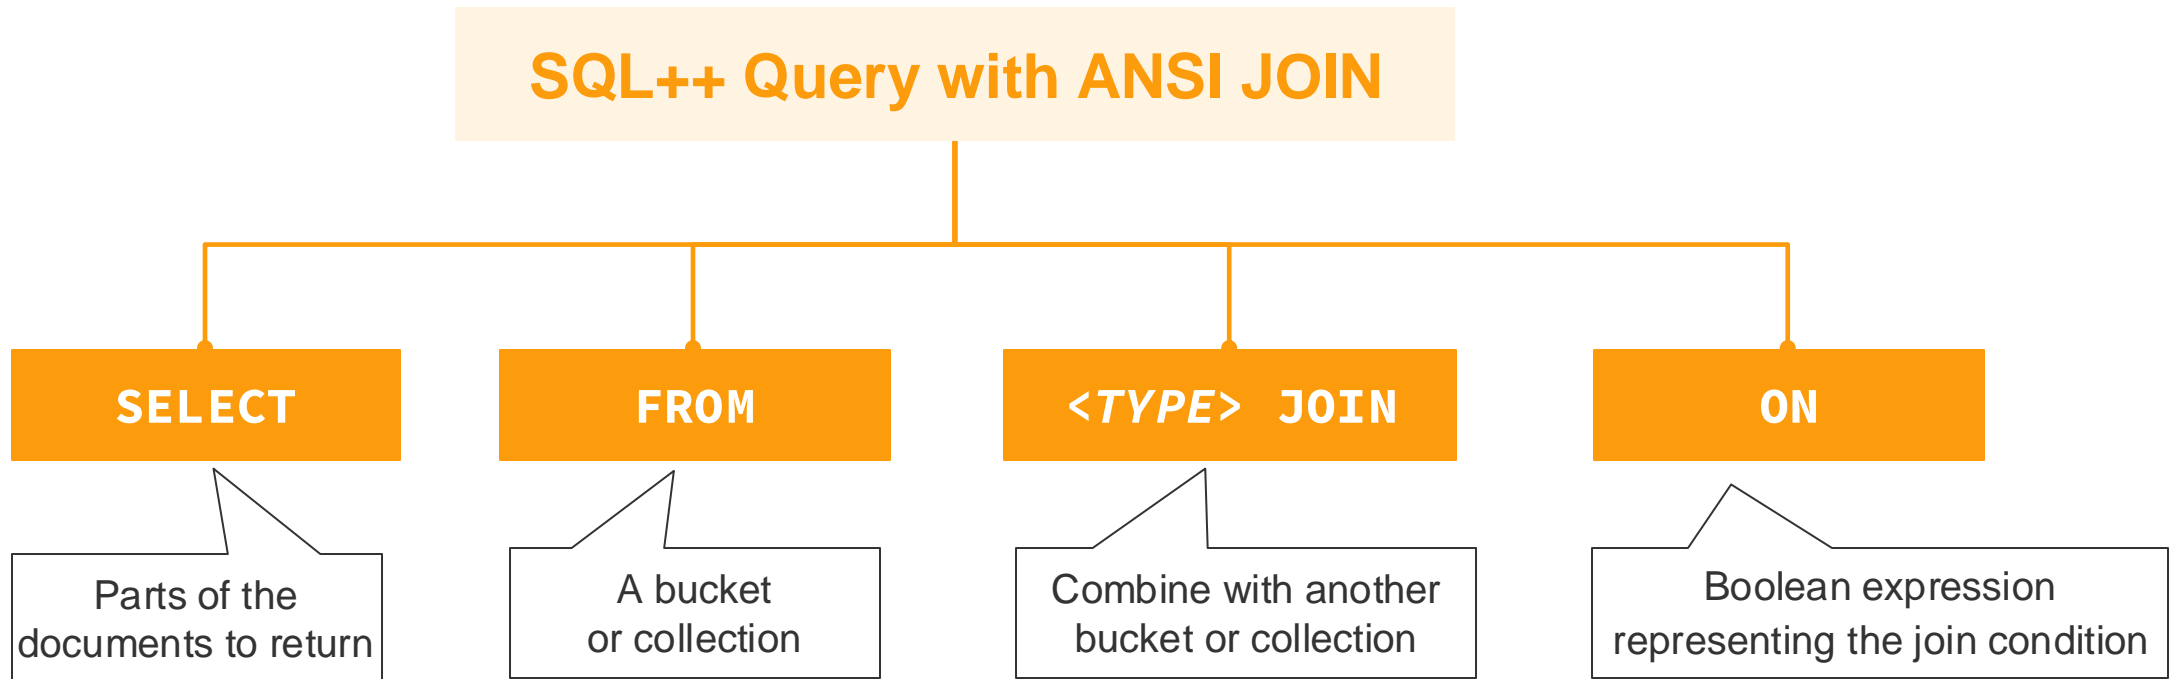
JOIN predicate

Update the routes using equipment 319

Delete routes having 757 equipment

- Provides the ability to modify a collection based on the results of a join with another collection.
- Based on a match or no match in the join, actions can be INSERT, UPDATE, DELETE.
- Multiple actions can be specified.

# SQL-like queries | ANSI JOINS



# SQL-like queries | ANSI JOINS

JOIN TYPE	Examples
<b>ANSI JOIN</b>	<pre>SELECT .. FROM customer c JOIN orders o ON o.customer_id = c.id</pre>
<b>ANSI JOIN Complex</b>	<pre>SELECT .. FROM airline JOIN route ON route.airlineid = "airline_"    toString(airline.id) AND route.type = "route"</pre>
<b>ANSI JOIN with IN CLAUSE</b>	<pre>SELECT .. FROM `travel-sample` route JOIN airport ON airport.faa IN [ route.sourceairport, route.destinationairport ] AND airport.type = "airport"</pre>
<b>ANSI LEFT OUTER</b>	<pre>SELECT .. FROM airport LEFT JOIN route ON airport.faa = route.sourceairport AND route.type = "route"</pre>
<b>ANSI JOIN with HASH JOIN</b>	<pre>SELECT .. FROM airport JOIN route USE HASH(build) ON airport.faa = route.sourceairport AND route.type = "route"</pre>

ANSI JOIN can join arbitrary fields of the documents and can be chained together.



# SQL-like queries | Arrays

*Retrieve the details of KL flight schedules from Albuquerque (ABQ) to Atlanta (ATL) if any of the flights are after 23:40.*

```
SELECT *
FROM `travel-sample`.inventory.route
WHERE airline="KL"
      AND sourceairport="ABQ"
      AND destinationairport="ATL"
      AND ANY departure IN schedule
          SATISFIES departure.utc > "23:40" END;
```



```
[
  {
    "travel-sample": {
      "airline": "KL",
      "airlineid": "airline_3090",
      "destinationairport": "ATL",
      "distance": 2038.3535078909663,
      "equipment": "757 320",
      "id": 36159,
      "schedule": [
        {"day": 0, "flight": "KL938", "utc": "03:54:00"},
        // ...
        {"day": 5, "flight": "KL169", "utc": "23:41:00"},
        // ...
        {"day": 6, "flight": "KL636", "utc": "17:40:00"}
      ],
      "sourceairport": "ABQ",
      "stops": 0,
      "type": "route"
    }
  }
]
```

- Range predicates (ANY, EVERY) enable you to evaluate expressions over every element in an array
- They are particularly useful when used to evaluate expressions over an array of objects

# SQL-like queries | Subqueries

*Find total number of airports by country where each city has more than 5 airports.*

```
SELECT t1.country,  
       array_agg(t1.city),  
       sum(t1.city_cnt) as apnum  
FROM  
  (SELECT city,  
         city_cnt,  
         array_agg(airportname) as apnames,  
         country  
   FROM `travel-sample`.inventory.airport  
  GROUP BY city,  
         country LETTING city_cnt = count(city)  
 ) AS t1  
WHERE t1.city_cnt > 5  
GROUP BY t1.country;
```



```
[  
  {  
    "$1": [  
      "Paris"  
    ],  
    "apnum": 9,  
    "country": "France"  
  },  
  {  
    "$1": [  
      "London"  
    ],  
    "apnum": 13,  
    "country": "United Kingdom"  
  },  
  {  
    "$1": [  
      "Houston",  
      "New York",  
      "San Diego"  
    ],  
    "apnum": 22,  
    "country": "United States"  
  }  
]
```

- A subquery is a query within another query
- Subqueries can be embedded anywhere a valid expression can go

# SQL-like queries | Built-in Functions

Category	Functions
Aggregate functions	ARRAY_AGG() ARRAY_AGG(DISTINCT) AVG(), AVG(DISTINCT) COUNT() COUNT(DISTINCT) MAX() MIN() SUM() SUM(DISTINCT)
Object functions	OBJECT_LENGTH() OBJECT_NAMES() OBJECT_PAIRS() OBJECT_INNER_PAIRS() OBJECT_VALUES() OBJECT_INNER_VALUES() OBJECT_ADD() OBJECT_PUT() OBJECT_REMOVE() OBJECT_UNWRAP()
Array functions	ARRAY_APPEND() ARRAY_AVG() ARRAY_CONCAT() ARRAY_CONTAINS() ARRAY_COUNT() ARRAY_DISTINCT() ARRAY_IFNULL() ARRAY_LENGTH() ARRAY_MAX() ARRAY_MIN() ARRAY_POSITION() ARRAY_PREPEND() ARRAY_PUT() ARRAY_RANGE() ARRAY_REMOVE() ARRAY_REPEAT() ARRAY_REPLACE() ARRAY_REVERSE() ARRAY_SORT() ARRAY_SUM()
Comparison functions	GREATEST() LEAST()
Conditional functions	IFMISSING() IFMISSINGORNULL() IFNULL() MISSINGIF() NULLIF() IFINF() IFNAN() IFNANORINF() NANIF() NEGINFIF() POSINFIF()
Number functions	ABS() ACOS() ASIN() ATAN() ATAN2() CEIL() COS() DEGREES() E() EXP() LN() LOG() FLOOR() PI() POWER() RADIANS() RANDOM() ROUND() SIGN() SIN() SQRT() TAN() TRUNC()
Date functions	CLOCK_MILLIS() CLOCK_STR() DATE_ADD_MILLIS() DATE_ADD_STR() DATE_DIFF_MILLIS() DATE_DIFF_STR() DATE_PART_MILLIS() DATE_PART_STR() DATE_TRUNC_MILLIS() DATE_TRUNC_STR() STR_TO_MILLIS() MILLIS_TO_STR() MILLIS_TO_UTC() MILLIS_TO_ZONE_NAME() NOW_MILLIS() NOW_STR() STR_TO_MILLIS() STR_TO_UTC() STR_TO_ZONE_NAME()
JSON functions	DECODE_JSON() ENCODE_JSON() ENCODED_SIZE() POLY_LENGTH()
Meta and UUID functions	BASE64() BASE64_ENCODE() BASE64_DECODE() META() UUID()
Pattern-matching functions	REG_CONTAINS() REG_LIKE() REG_POSITION() REG_REPLACE()
String functions	CONTAINS() INITCAP() LENGTH() LOWER() LTRIM() POSITION() REPEAT() REPLACE() RTRIM() SPLIT() SUBSTR() TITLE() TRIM() UPPER()
Type-Checking Functions	ISARRAY() ISATOM() ISBOOLEAN() ISNUMBER() ISOBJECT() ISSTRING() TYPE()
Type-Conversion Functions	TOARRAY() TOATOM() TOBOOLEAN() TONUMBER() TOOBJECT() TOSTRING()

# SQL-like queries | User Defined Functions (UDF)

## Simple Inline Example

```
CREATE FUNCTION to_meters(...) {  
  args[0] * 0.3048  
};
```

```
SELECT airportname, ROUND(to_meters(geo.alt)) AS mamsl  
FROM `travel-sample`.inventory.airport  
LIMIT 5;
```



```
[  
  {"airportname": "Calais Dunkerque", "mamsl": 4},  
  {"airportname": "Peronne St Quentin", "mamsl": 90},  
  {"airportname": "Les Loges", "mamsl": 130},  
  {"airportname": "Couterne", "mamsl": 219},  
  {"airportname": "Bray", "mamsl": 111}  
]
```

## Inline Example with SQL++ query inside

```
CREATE FUNCTION locations(vActivity) {  
  (SELECT id, name, address, city  
   FROM `travel-sample`.inventory.landmark  
   WHERE activity = vActivity) };
```

```
SELECT l.name, l.city  
FROM locations("eat") AS l  
WHERE l.city = "Gillingham";
```



```
[  
  {"city": "Gillingham", "name": "Hollywood Bowl"},  
  {"city": "Gillingham", "name": "Thai Won Mien"},  
  {"city": "Gillingham", "name": "Spice Court"},  
  {"city": "Gillingham", "name": "Beijing Inn"},  
  {"city": "Gillingham", "name": "Ossie's Fish and  
  Chips"}  
]
```

- Inline functions are defined using SQL++ expressions, including subqueries.
- You can name and reuse complex or repetitive expressions in order to simplify your queries.

# SQL-like queries | UDF with JavaScript & SQL++

Create a Javascript function in a Library

demodml

```
1 /* a UDF library contains one or more javascript functions */
2 = function getairportname(airportcode)
3   var query = SELECT a.airportname FROM `travel-sample`.inventory.airport
4   a WHERE a.icao=$airportcode;
5   let acc = [];
6   for (const row of query) {
7     acc.push(row);
8   }
9
10  return(acc);
11 }
```

Javascript library

Create a UDF Function that references it

Edit Function

Function Name  
getname

Namespace  
global

Parameters  
--

Function Type  
javascript

Javascript Library  
demodml

Library Function Name  
getairportname

Use the UDF function in SQL++ statements

Query Editor

1 select getname(a.icao) from `travel-sample`.inventory.airport a limit 1;

Execute Run as TX Index Advisor Explain

success 2 min ago 247.5ms | 1 docs | 111 bytes

Results

Table JSON Chart Plan Plan Text Advice

1- [
2- {
3- "\$1": [
4- {
5- "airportname": "Calais Dunkerque"
6- }
7- ]
}

- UDFs with JavaScript allows developers to provide custom functions to extend SQL++ capabilities

# SQL-like queries | Search Functions

*Find the name of the hotels in United Kingdom where the reviews match the term 'bathrobes'*

```
SELECT t1.name, meta().id
FROM `travel-sample`.inventory.hotel AS t1
WHERE SEARCH(t1, {
  "match": "bathrobes",
  "field": "reviews.content",
  "analyzer": "standard"
})
AND country="United Kingdom"
LIMIT 3;
```



```
[
  {
    "id": "hotel_12068",
    "name": "Castle Hotel"
  },
  {
    "id": "hotel_18819",
    "name": "Bistro Prego With Rooms"
  },
  {
    "id": "hotel_3622",
    "name": "Premier Inn Birmingham Central East"
  }
]
```

- Search functions enable you to use full text search (FTS) queries directly within a SQL++ query.
- It is recommended that you create suitable full text indexes for the searches that you need to perform.

# SQL-like queries | Prepared Statement

## Positional Parameters

```
PREPARE NumParam AS  
SELECT * FROM `travel-sample`.inventory.hotel  
WHERE city=$1 AND country=$2;
```

```
EXECUTE NumParam  
USING ["Paris", "France"];
```

## Named Parameters

```
PREPARE NameParam AS  
SELECT * FROM `travel-sample`.inventory.hotel  
WHERE city=$city AND country=$country;
```

```
EXECUTE NameParam  
USING {"city": "Paris", "country": "France"};
```

- You can add placeholder parameters to a statement, so that you can supply variable values when you run the statement.
- If you need to run a statement more than once, you can prepare the execution plan for the statement.

# SQL-like queries | Transactions

Ensures database consistency when multiple documents are updated in a single or multiple SQL++ statements

## **START TRANSACTION;**

```
UPDATE customer SET balance = balance + 100 WHERE cid = 4872;  
SELECT cid, name, balance FROM customer;
```

## **SAVEPOINT s1;**

```
UPDATE customer SET balance = balance - 100 WHERE cid = 1924;  
SELECT cid, name, balance FROM customer;
```

## **ROLLBACK WORK TO SAVEPOINT s1;**

```
SELECT cid, name, balance FROM customer;
```

## **COMMIT;**

Couchbase provides the following statements for transactions:

- BEGIN TRANSACTION
- SET TRANSACTION (optional)
- SAVEPOINT
- ROLLBACK TRANSACTION
- COMMIT TRANSACTION





# 수고하셨습니다.



[paul.son@couchbase.com](mailto:paul.son@couchbase.com)

[www.couchbase.com](http://www.couchbase.com)

[cloud.couchbase.com](http://cloud.couchbase.com)



**Couchbase**